

**KTH**

DEGREE PROJECT IN COMPUTER SCIENCE AND ENGINEERING  
SECOND CYCLE, 30 CREDITS

Stockholm, Sweden 2026

# Joint Radar Emitter and Mode Clustering Using a Transformer Encoder Architecture

Markus Swegmark

Supervisor: TBD

Examiner: TBD

Host: TBD

KTH Royal Institute of Technology  
School of Electrical Engineering and Computer Science  
Master's Programme in Machine Learning

TRITA-EECS-EX-XXXX:XXX

# Abstract

This thesis investigates joint clustering of radar emitters and their operating modes from intercepted pulse descriptor word (PDW) streams using a transformer encoder architecture.

**Keywords:** transformer, deep clustering, electronic intelligence, radar emitter identification, mode recognition, supervised contrastive learning.

# Sammanfattning

Detta examensarbete undersöker gemensam klustering av radarsändare och deras operativa moder från avlyssnade pulsbeskrivande ord (PDW) med hjälp av en transformerbaserad encoderarkitektur.

**Nyckelord:** transformer, djup klustering, signalspaning, radarsändaridentifiering, modigenkänning, självövervakad inlärning.

# Acknowledgments

I would like to thank my supervisor and examiner at KTH, and the host organization for their support throughout this thesis project.

*Stockholm, June 2026*

Markus Swegmark

# Contents

|                                                                        |            |
|------------------------------------------------------------------------|------------|
| <b>Abstract</b>                                                        | <b>i</b>   |
| <b>Sammanfattning</b>                                                  | <b>ii</b>  |
| <b>Acknowledgments</b>                                                 | <b>iii</b> |
| <b>List of Acronyms</b>                                                | <b>x</b>   |
| <b>1 Introduction</b>                                                  | <b>1</b>   |
| 1.1 Background and Motivation . . . . .                                | 1          |
| 1.2 Problem Statement . . . . .                                        | 1          |
| 1.3 Research Questions . . . . .                                       | 2          |
| 1.4 Objectives and Scope . . . . .                                     | 2          |
| 1.4.1 Objectives . . . . .                                             | 2          |
| 1.4.2 Delimitations . . . . .                                          | 3          |
| 1.5 Contributions . . . . .                                            | 3          |
| 1.6 Ethical and Societal Considerations . . . . .                      | 3          |
| 1.7 Thesis Outline . . . . .                                           | 4          |
| <b>2 Background</b>                                                    | <b>5</b>   |
| 2.1 Radar Systems and Passive Reception . . . . .                      | 5          |
| 2.1.1 Electromagnetic Pulses for Navigation, Surveillance, and Control | 5          |
| 2.1.2 Time-of-Flight Ranging and Ambiguity Fundamentals . . . . .      | 5          |
| 2.1.3 Transmitting radars versus receive-only systems . . . . .        | 6          |
| 2.1.4 From the raw signal to a pulse sequence and descriptor words .   | 7          |
| 2.1.5 Physical sources and transmit schedules . . . . .                | 8          |
| 2.2 Signal Processing Pipeline for Passive ELINT . . . . .             | 9          |
| 2.2.1 Front-end reception, detection, and pulse-parameter estimation . | 9          |
| 2.2.2 Deinterleaving and pulse-train formation . . . . .               | 9          |
| 2.2.3 Emitter reporting and equipment identification . . . . .         | 10         |
| 2.2.4 Operating-mode analysis and waveform libraries . . . . .         | 10         |
| 2.2.5 Emitter geolocation, fusion, and downstream displays . . . . .   | 11         |
| 2.3 Machine Learning Fundamentals . . . . .                            | 11         |
| 2.3.1 Learning paradigms . . . . .                                     | 11         |
| 2.3.2 Representation learning . . . . .                                | 12         |
| 2.3.3 Applications in radar deinterleaving and mode classification . . | 12         |
| 2.4 Clustering Methods . . . . .                                       | 13         |
| 2.5 Deep Representation Learning . . . . .                             | 13         |

|          |                                                             |           |
|----------|-------------------------------------------------------------|-----------|
| 2.6      | The Transformer Encoder . . . . .                           | 14        |
| <b>3</b> | <b>Related Work</b>                                         | <b>16</b> |
| 3.1      | Classical Radar Emitter Identification . . . . .            | 16        |
| 3.2      | Deep Learning for Radar Signal Analysis . . . . .           | 16        |
| 3.3      | Deep Clustering and Joint Representation Learning . . . . . | 16        |
| 3.4      | Transformers for Time Series and Sets . . . . .             | 17        |
| 3.5      | Joint and Multi-Task Clustering . . . . .                   | 17        |
| <b>4</b> | <b>Method</b>                                               | <b>18</b> |
| 4.1      | Problem Formulation . . . . .                               | 18        |
| 4.2      | Data Representation and Preprocessing . . . . .             | 20        |
| 4.2.1    | Windowing . . . . .                                         | 20        |
| 4.2.2    | Per-feature scaling . . . . .                               | 21        |
| 4.2.3    | Labels . . . . .                                            | 22        |
| 4.3      | Transformer Encoder Architecture . . . . .                  | 22        |
| 4.4      | Loss Function . . . . .                                     | 26        |
| 4.5      | Training Procedure . . . . .                                | 28        |
| 4.6      | Evaluation Metrics . . . . .                                | 29        |
| 4.7      | Baseline Methods . . . . .                                  | 33        |
| <b>5</b> | <b>Experiments and Results</b>                              | <b>34</b> |
| 5.1      | Experimental Setup . . . . .                                | 34        |
| 5.2      | Datasets . . . . .                                          | 34        |
| 5.2.1    | Qualitative properties of the synthetic corpus . . . . .    | 35        |
| 5.3      | Hyperparameters and Training Details . . . . .              | 36        |
| 5.4      | Quantitative Results . . . . .                              | 36        |
| 5.5      | Ablation Studies . . . . .                                  | 37        |
| 5.6      | Qualitative Analysis . . . . .                              | 37        |
| <b>6</b> | <b>Discussion</b>                                           | <b>38</b> |
| 6.1      | Interpretation of Results . . . . .                         | 38        |
| 6.2      | Comparison with Related Work . . . . .                      | 38        |
| 6.3      | Limitations . . . . .                                       | 39        |
| 6.4      | Implications . . . . .                                      | 40        |
| <b>7</b> | <b>Conclusion</b>                                           | <b>41</b> |
| 7.1      | Summary . . . . .                                           | 41        |
| 7.2      | Answers to the Research Questions . . . . .                 | 41        |
| 7.3      | Future Work . . . . .                                       | 41        |
|          | <b>References</b>                                           | <b>42</b> |

---

|                                 |           |
|---------------------------------|-----------|
| <b>A Additional Results</b>     | <b>44</b> |
| <b>B Implementation Details</b> | <b>45</b> |
| <b>C Notation Reference</b>     | <b>46</b> |

# List of Figures

|     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |    |
|-----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 2.1 | Monostatic radar geometry and round-trip timing. A co-located transmit/receive (TX/RX) aperture on the left emits a pulse that travels to a target at range $R$ , scatters, and returns to the same aperture as an echo. The lower timeline shows the corresponding events: the outgoing TX pulse at $t = 0$ and the received echo at $t = \tau$ , separated by the round-trip delay $\tau = 2R/c$ from Equation (2.1). . . . .                                                                                                                                                                                                                                                | 6  |
| 2.2 | Schematic of the parameters captured by a PDW for a short pulse train. Each pulse contributes a TOA (the leading edge $t_n$ on the time axis), a pulse width (PW), a peak amplitude $A_n$ , and a carrier RF $f_n$ , sketched as the fast oscillation inside the middle pulse; successive pulses also define a PRI (equivalently a $\text{PRF} = 1/\text{PRI}$ ). The inset shows the AOA $\theta_n$ measured at the receive aperture relative to a boresight reference. A typical PDW stacks these per-pulse measurements into a vector $\mathbf{x}_n = (f_n, \text{PW}_n, t_n, A_n, \theta_n)$ that is the usual input to downstream emitter and mode reasoning [1]. . . . . | 8  |
| 4.1 | The two label-induced partitions of a single pulse train. Each pulse $\mathbf{x}_n$ carries two labels: an emitter identity $e_n$ (encoded as fill colour) and an operating mode $m_n$ (encoded as border style: solid, dashed, or dotted). The same index set $\{1, \dots, N\}$ is partitioned in two distinct ways, by the emitter labelling (Equation (4.1)) and by the mode labelling (Equation (4.2)). Within each emitter subset the modes are in general mixed, and within each mode subset the emitters are in general mixed, reflecting the coupling between the two labellings. . . . .                                                                              | 19 |
| 4.2 | Sliding-window segmentation of a pulse train as defined in Equation (4.6). The construction is valid for any stride $0 < s \leq L$ ; the figure uses $s = L/2$ for legibility. Successive windows $\mathbf{X}^{(w)}$ advance by $s$ pulses and therefore overlap by $L - s$ . Dashed guides project window boundaries down onto the pulse train. Pulses past the right edge of the last window of length $L$ form a trailing remainder of length $< L$ that is dropped during training and evaluation. Training uses $s = L/2$ and validation/test use $s = L$ , the latter giving abutting non-overlapping windows. . . . .                                                   | 21 |



|     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |    |
|-----|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 4.3 | Architecture of the encoder $f_{\theta}$ . A sequence of PDWs is first lifted to $d_{\text{model}} = 256$ by an input embedding, processed by a shared trunk of 6 transformer-encoder layers, and then split into two parallel branches with their own $256 \rightarrow 128$ projection and 2 encoder layers. Branch outputs are mapped by linear projection heads to the emitter embedding space $\mathbb{R}^8$ and the mode embedding space $\mathbb{R}^{16}$ . . . . . | 23 |
| 4.4 | Inference-time decoding of branch embeddings into per-pulse cluster indices via Equation (4.14). Dots are pulse embeddings and dashed ellipses are DBSCAN clusters in one window. The pipeline is identical on both branches and the integer outputs are local to each clustering instance; they acquire their interpretation from the supervision scope of Section 4.1, which differs between the two branches. . . . .                                                  | 25 |

# List of Tables

|     |                                                                                                                                             |    |
|-----|---------------------------------------------------------------------------------------------------------------------------------------------|----|
| 5.1 | Synthetic scenario corpus used for training and evaluation. Dashes mark quantities still being reconciled with the export manifest. . . . . | 35 |
| 5.2 | Primary optimisation hyperparameters for the proposed model. . . . .                                                                        | 36 |

# List of Acronyms

**ACC** clustering accuracy

**AMI** adjusted mutual information

**AOA** angle of arrival

**ARI** adjusted Rand index

**BERT** Bidirectional Encoder Representations from Transformers

**DBSCAN** Density-based spatial clustering of applications with noise

**DEC** deep embedded clustering

**DL** deep learning

**ELINT** electronic intelligence

**ESM** electronic support measures

**EW** electronic warfare

**FFN** feed-forward network

**HDBSCAN** Hierarchical density-based spatial clustering of applications with noise

**MHA** multi-head attention

**ML** machine learning

**MLP** multilayer perceptron

**NMI** normalized mutual information

**PDW** pulse descriptor word

**PRF** pulse repetition frequency

**PRI** pulse repetition interval

**RADAR** radio detection and ranging

**RF** radio frequency

**TOA** time of arrival

**V-measure** harmonic mean of homogeneity and completeness

# CHAPTER 1

## Introduction

### 1.1 Background and Motivation

Recent large-scale conflicts illustrate how contested the electromagnetic spectrum has become. Commanders rely on timely, passive sensing to build situational awareness, protect platforms, and coordinate fires without revealing their own emissions. Within EW, which aims to exploit, protect, and maneuver in the spectrum, ESM systems listen passively and supply the measurements that feed higher-level ELINT tasks such as detecting radars, grouping pulses by emitter, and inferring behavior [1, 2].

Modern radars are increasingly agile: they vary carrier frequency, pulse width, amplitude, and PRI across bursts to improve performance and survivability. That agility blurs simple fingerprints and stresses classical rule-based parsers tuned to stable pulse trains. The same physical emitter also switches *modes*—recurrent control policies that map an operational objective (for example, volume search versus track) to a characteristic regime in the pulse train. Mode awareness is therefore operationally informative: it constrains how the waveform may evolve, signals intent at a coarse level, and can help disambiguate emitters when type-level signatures overlap.

At the same time, ML and DL have become the default toolkit for sequence modeling and representation learning in other domains. Radar pulse analysis naturally presents as a stream of low-level measurements—often summarized as PDWs—where long-range dependence and mixture structure (multiple emitters superposed in time) favor learned encoders over purely hand-crafted feature pipelines, and where the available supervision is itself asymmetric: emitter identifiers in curated training data are typically defined per pulse train rather than across an open global population, while operating modes are drawn from a much smaller, shared catalogue.

This thesis studies whether a transformer-style encoder, trained with objectives suited to clustering, can produce embeddings in which both distinct emitters and their operating modes emerge as separable structure. The remainder of the introduction states the problem precisely, lists the research questions, and outlines scope, contributions, and ethics.

### 1.2 Problem Statement

Passive ESM produces a time-ordered stream of PDWs in which pulses from multiple emitters are interleaved and each emitter may change operating mode over time (Section 1.1). Operators ultimately require both emitter grouping and operating-mode

information from the same measurement stream, and classical pipelines therefore chain deinterleaving with separate mode heuristics or specialised supervised classifiers; under agile waveforms and overlapping trains, that separation scales poorly and does not guarantee a single representation whose geometry exposes emitter- and mode-level structure simultaneously. This thesis addresses the joint problem in a supervised setting: training pulse trains are generated by a controlled scenario simulator (Section 4.2) in which every PDW carries two ground-truth labels with asymmetric scope — a per-pulse emitter identity that is unique only *within* the pulse train that contains the pulse, and an operating mode drawn from a finite catalogue that is shared across all pulse trains. The task is therefore to learn from this label structure a single representation in which (i) the train-local emitter labels can be used to cluster pulses by physical emitter at inference time, including on previously unseen trains whose number of emitters is not known in advance, and (ii) the global mode labels can be used to cluster pulses by operating mode, including on modulation types not present in the training catalogue, while remaining robust to mixture effects, noise, and imperfect preprocessing.

## 1.3 Research Questions

The thesis addresses the following research questions:

- RQ1** Under practical sequence-length limits of a transformer encoder, can the model support effective deinterleaving and recover operating modes via clustering?
- RQ2** Does a joint objective that couples deinterleaving with mode-aware clustering yield better clustering performance or representations than a comparable single-objective transformer baseline focused on deinterleaving alone?

## 1.4 Objectives and Scope

### 1.4.1 Objectives

The work pursues objectives that operationalize the research questions in Section 1.3. First, we specify and implement a transformer-style encoder that ingests bounded windows of PDWs from interleaved pulse streams and maps them to a dual embedding (Section 4.1), together with a post-hoc clustering stage applied independently to the emitter and the mode branches. We then study how performance depends on practical limits on context length, mixture density, and noise, reporting agreement with the ground-truth simulator partitions using standard scores such as NMI and ARI (Chapter 5). Second, we formulate at least one joint training objective that couples a within-train supervised contrastive term for emitter identity with a batch-wide supervised contrastive term over the global mode labels, and compare it against a matched baseline trained primarily for deinterleaving or representation quality without

that coupling, holding architecture capacity and data conditions as constant as feasible (Chapters 4 and 5).

### 1.4.2 Delimitations

The scope is deliberately narrow. We operate at the level of PDWs or equivalent tabular pulse features supplied to the model; pulse detection, RF conditioning, antenna-level processing, and geometry-driven fusion across platforms are out of scope. The focus is offline or batch analysis of recorded streams rather than hard real-time implementation on embedded hardware, and we do not address full operational emitter naming against classified libraries. That choice excludes an *online* regime in which models update continuously from a live pulse stream under latency and memory budgets: training and evaluation here assume access to bounded windows that can be revisited, shuffled, and mini-batched, which sidesteps causal filtering, drift handling, and resource accounting that a streaming deployment would require. Consequently, conclusions concern representation quality and clustering agreement on curated windows rather than end-to-end tracking performance over arbitrarily long horizons. Training relies on the simulator-provided emitter and mode labels (Section 4.2); at inference time the trained encoder is intended to operate on previously unseen pulse trains whose emitter and mode memberships are both recovered by clustering the corresponding embedding trajectories, as described in Section 1.2. Finally, we do not claim field validation on classified measurements; conclusions are drawn from synthetic or curated datasets together with stress tests that mimic deinterleaving noise (Chapter 5).

## 1.5 Contributions

The main contributions of this thesis are:

- Placeholder contribution.
- Placeholder contribution.

## 1.6 Ethical and Societal Considerations

Methods for passive ESM and ELINT sit squarely in dual-use territory: the same representations that support defensive situational awareness, spectrum management, and safety-of-navigation applications can, in principle, support targeting or surveillance against specific platforms if deployed without appropriate policy, oversight, and legal constraints. This thesis develops ML techniques on abstract pulse-level features and does not address deployment, rules of engagement, or classification of live operational data; nevertheless, the capability to group emitters and recognise operating modes from

PDW streams should be understood as a general-purpose building block whose societal consequences depend on the institutional context in which it is used.

From a data-ethics perspective, the experimental work relies on controlled or synthetic PDW streams rather than collections derived from live emitters, which limits privacy and classification sensitivity concerns at the cost of weaker guarantees about behavior under classified or vendor-specific waveforms. When such methods are extended toward operational data, stewards must enforce access control, purpose limitation, and retention policies consistent with national regulation and organizational policy, and they should document provenance and limitations so that operators do not over-trust cluster labels that lack ground truth.

Training modern DL encoders incurs non-negligible energy use for computation and cooling; relative to lightweight classical deinterleaving heuristics, transformer-scale models increase the carbon footprint of experimentation and of any future always-on training loops. Mitigation follows standard practice: reuse checkpoints, avoid redundant hyperparameter sweeps, prefer smaller models when they meet accuracy targets, and report approximate training cost where it informs reproducibility.

At a societal level, automation of emitter and mode analysis may concentrate analytic capacity in well-resourced actors, widen asymmetries in electromagnetic situational awareness, and reduce the visibility of human expert judgment if clusters are treated as authoritative without audit trails. Responsible research therefore favors transparent evaluation protocols, clear statements of failure modes under noise and mixture, and conservative claims about operational readiness.

## 1.7 Thesis Outline

The remainder of the thesis is organized as follows. Chapter 2 supplies textbook-style background on pulsed radar, PDWs, emitters and operating modes, the ELINT processing chain, and the ML and transformer concepts needed to read the method. Chapter 3 turns to the literature, contrasting classical emitter identification with deep supervised radar models, deep clustering objectives, sequence transformers, and work on joint or hierarchical cluster structure; each subsection closes with the limitation that motivates the present approach. Chapter 4 states the formal problem, documents data representation and preprocessing (Section 4.2), specifies the dual-branch encoder and losses (Sections 4.3 and 4.4), describes training and metrics, and lists baselines. Chapter 5 reports hardware and software setup, the synthetic scenario corpus (Section 5.2), hyperparameters, quantitative scores, ablations, and qualitative analyses. Chapter 6 interprets findings, positions them against prior art, states limitations—including simulator-only evidence and operational gaps—and discusses implications. Chapter 7 summarizes answers to the research questions and outlines future work.

## CHAPTER 2

# Background

## 2.1 Radar Systems and Passive Reception

### 2.1.1 Electromagnetic Pulses for Navigation, Surveillance, and Control

The term RADAR stands for **R**adio **D**etection and **R**anging. At its core, a radar system operates by sending out electromagnetic pulses and then listening for the echoes that reflect back from objects in the environment. This basic principle supports a wide range of applications, including air traffic control, weather monitoring, altitude measurement, navigation of ships and vehicles, target tracking, and monitoring of the electromagnetic spectrum alongside commercial communication services [2].

What unifies all radar systems is a practical goal: they emit a carefully controlled burst of electromagnetic energy, allow it to scatter from surrounding objects, and then extract information about the location and motion of those objects from the returning signals. Electromagnetic waves are generated whenever electric charges accelerate, such as when an alternating current passes through an antenna. These waves move outward from the antenna at the speed of light in air. In radar, each outgoing pulse can be thought of as a brief and well-defined packet of energy. By recording the time it takes for this energy to bounce off an object and return, and knowing the speed at which it traveled, the radar can directly calculate how far away the object is [2].

### 2.1.2 Time-of-Flight Ranging and Ambiguity Fundamentals


At its core, radar sensing uses the physics of wave propagation to measure distance. A transmitter emits a brief electromagnetic pulse, which travels at speed  $c$  (the speed of light in air), reflects from an object, and returns as an echo. By measuring the elapsed time  $\tau$  between emission and reception—known as the round-trip delay—the system can infer the distance to the object.

In the canonical configuration, a single (monostatic) radar platform handles both transmission and reception, with the wave traversing to the target and back. The fundamental relation connecting time delay to range follows directly from the definition of speed:

$$R = \frac{c\tau}{2}, \tag{2.1}$$

where  $R$  is the target range and  $c$  is the wave speed [2]. The factor of 2 accounts for the outbound and return paths.





**Figure 2.1.** Monostatic radar geometry and round-trip timing. A co-located transmit/receive (TX/RX) aperture on the left emits a pulse that travels to a target at range  $R$ , scatters, and returns to the same aperture as an echo. The lower timeline shows the corresponding events: the outgoing TX pulse at  $t = 0$  and the received echo at  $t = \tau$ , separated by the round-trip delay  $\tau = 2R/c$  from Equation (2.1).

A similar approach is found in animal echolocation: for instance, a bat produces a short sound pulse and estimates the distance to an obstacle or prey by interpreting the delay before it hears the echo, substituting sound waves for electromagnetic ones [2].

While monostatic setups are common, many practical radar and electronic support systems employ multiple spatially separated receive antennas, or distributed arrays, forming so-called bistatic or multistatic configurations. By correlating the signals arriving at different locations, these systems can measure angles of arrival and other spatial properties in addition to range, thus supporting direction finding and improved localization.

Practical radar systems emit periodic pulses, each separated by a pulse repetition interval (PRI) or, equivalently, a pulse repetition frequency (PRF). A core challenge appears when echoes from earlier pulses arrive after a new pulse is sent. This leads to *range ambiguity*, where it is unclear which transmitted pulse a received echo should be attributed to. The largest range without ambiguity—the *unambiguous range*—is set by the time between pulses:

$$R_u = \frac{c}{2 \text{PRF}} = \frac{c \text{PRI}}{2}. \quad (2.2)$$

Signal processing techniques such as matched filtering and coherent integration increase detection performance within this range, but no algorithm can resolve targets unambiguously beyond  $R_u$  as limited by the physics of pulse timing [2].

### 2.1.3 Transmitting radars versus receive-only systems

Any receiver within range can detect intentional electromagnetic transmissions, and repeated radar pulses make timing and direction information available not only to the radar user, but to anyone listening in [1]. This exposes a fundamental trade-off: transmitting enables the measurement of objects at a distance, but simultaneously reveals the radar’s presence, position, and activity to others in the environment.

Traditional (transmitting) radar systems measure the delay between their own outgoing pulses and the returning echoes, allowing direct range estimation [2]. In contrast, passive systems such as ESM and ELINT do not transmit their own signals. Instead, they monitor emissions from other sources and attempt to infer information about those emitters and their behavior [1]. Because they do not transmit, passive receivers maintain a lower profile and avoid revealing themselves by radiation—a key advantage in situations where stealth is important.

Lacking direct control over the transmitted waveform, passive receivers rely on whatever pulse measurements they can extract after detection: carrier frequency, time of arrival, pulse width, amplitude, and angle of arrival if available. This necessity motivates the PDW abstraction described next. Section 2.2 shows how PDWs fit into the overall signal processing chain.

### 2.1.4 From the raw signal to a pulse sequence and descriptor words

Every radar or receiver starts by collecting a raw electrical signal—essentially a voltage that varies in time as electromagnetic waves are picked up by the antenna. This raw signal can be viewed as a digital recording of the environment’s radio activity, capturing everything present, including both the pulses of interest and background noise.


To extract information, the system looks for brief increases in signal energy that stand out above the noise. These increases correspond to pulses—short bursts of transmitted energy. Locating a pulse means finding where in time the signal rises above a set threshold. After locating pulses, the receiver estimates several physical properties for each one: when it arrived (its time of arrival), how long it lasted (its pulse width), how strong it was (amplitude or energy), and what frequency it had. If the receiver has directional antennas, it can also estimate the angle from which the pulse arrived.

After initial detection and measurement, each pulse is summarized into a standardized set of numbers called a *PDW* (pulse descriptor word). This record captures the essential measured properties of a pulse, such as:

$$\mathbf{x}_n = (f_n, \text{PW}_n, t_n, A_n, \theta_n, \dots)^\top \in \mathbb{R}^d, \quad (2.3)$$

where  $f_n$  is the measured carrier frequency,  $\text{PW}_n$  is pulse width,  $t_n$  is time of arrival (often the leading edge),  $A_n$  is amplitude, and  $\theta_n$  is angle of arrival if available. The ellipsis allows for extra recorded properties, like elevation angle of arrival, but the main idea is that each column represents a specific physical measurement, and the format is fixed once chosen. Figure 2.2 visualizes this five-parameter template.

A sequence of such pulses, indexed by order of arrival  $(\mathbf{x}_1, \mathbf{x}_2, \dots)$ , forms what is called a *pulse train*. For pulses that come from a single transmitter operating in a stable mode, the intervals between arrivals ( $\tau_n = t_n - t_{n-1}$ , for  $n \geq 2$ ) are usually regular or follow a set schedule. This interval is called the pulse repetition interval (PRI); its



**Figure 2.2.** Schematic of the parameters captured by a PDW for a short pulse train. Each pulse contributes a TOA (the leading edge  $t_n$  on the time axis), a pulse width (PW), a peak amplitude  $A_n$ , and a carrier RF  $f_n$ , sketched as the fast oscillation inside the middle pulse; successive pulses also define a PRI (equivalently a  $PRF = 1/PRI$ ). The inset shows the AOA  $\theta_n$  measured at the receive aperture relative to a boresight reference. A typical PDW stacks these per-pulse measurements into a vector  $\mathbf{x}_n = (f_n, PW_n, t_n, A_n, \theta_n)$  that is the usual input to downstream emitter and mode reasoning [1].

inverse, the pulse repetition frequency (PRF), describes how often pulses are sent:

$$\tau_n = t_n - t_{n-1}, \quad n \geq 2. \quad (2.4)$$

While simple systems use a fixed interval, more complex transmitters may change their intervals in a predictable or agile way, and this information can be included in the record for each pulse.

Each measured descriptor is thus a function of both the true (but unknown) properties of the transmitting source and the effects of noise, measurement error, and possible missing data. Very weak pulses may be missed, and noise or spurious events may be incorrectly identified as pulses.

Finally, rather than carrying forward the full raw signal, the system represents each pulse only by its measured properties in the form of a time-ordered sequence of PDWs. This sequence  $(\mathbf{x}_1, \dots, \mathbf{x}_T)$ —sometimes called the pulse train—is the main input to later processing and learning algorithms. Each  $\mathbf{x}_n$  is a compact summary of one pulse, and the order reflects the timeline of arrivals, replacing the need for direct inspection of the original sampled voltage trace.

### 2.1.5 Physical sources and transmit schedules

In passive ESM and ELINT, a *radar emitter* means a physical transmitter whose pulses line up in time and frequency strongly enough to be treated as one source [1]. An *operating mode* is the instantaneous transmit recipe: carrier plan, pulse width, repetition pattern, beam motion, and any programmed switching rules among them [1, 2]. The same hardware cycles among modes as tasks change, so a timeline of PDWs mixes mode segments even when the physical box is fixed. Different platforms can also implement

the same functional mode with different numerical parameters, for example two trackers with distinct nominal PRFs [1].

Illustrative mode families include wide-area search, acquisition and track, fire-control illumination, weather sensing, and terminal guidance support [2]. Those roles push different combinations of PRI, dwell, agility, and scan law, which is why operators care both about identity and about what the transmitter is doing at a given time [1, 2].

The relationship between emitters and modes is many-to-many: one emitter runs many modes over a mission, and different emitters can realize similar modes with different parameter sets [1]. The rest of this chapter walks through the passive processing chain (Section 2.2) and then supplies the learning background used in the method.

## 2.2 Signal Processing Pipeline for Passive ELINT

In EW and ESM applications, measured radar energy passes through a staged *signal processing pipeline* before it reaches higher-level ELINT reasoning. The exact implementation varies with platform, band, and mission software, but the logical chain is broadly stable across textbooks and system descriptions.

### 2.2.1 Front-end reception, detection, and pulse-parameter estimation

At the *front end*, antennas and RF conditioning capture emissions over one or more instantaneous bandwidths, often after wideband digitization or channelized sub-band processing. A *detection* stage then marks pulse-like events against a noise and clutter background, producing a sparse list of candidate pulse times (and sometimes coarse frequency hints).

Given detections, *parameter estimation* attaches per-pulse measurements such as carrier frequency, pulse width, time of arrival, direction of arrival when available, and amplitude or signal-to-noise proxies, this becomes our PDWs after possibly attaching AoA after triangulation. These estimates are noisy, can be missing for weak emitters, and may include false alarms that are not associated with any physical radar.

### 2.2.2 Deinterleaving and pulse-train formation

*Sorting* and *deinterleaving* group pulses into coherent pulse trains by exploiting periodic structure in inter-pulse timing and consistent parameter tracks across time. The output of this stage is commonly packaged as a stream of PDWs (or closely related pulse descriptor records), which form the usual interface to downstream tasks such as emitter identification and mode analysis (Section 2.1.4).

Operationally, deinterleaving developed from deterministic temporal rules toward multi-parameter clustering and, more recently, toward learned sequence models [3]. In

comparatively sparse environments with simple pulse trains, PRI inferred from TOA differences often behaved like a stable timing signature, so early algorithms searched for repetition by accumulating histograms of inter-pulse intervals. Cumulative and sequential difference histograms (CDIF and SDIF) are representative: they “vote” for candidate PRI values through peaks in the difference domain and remain standard pedagogical references [1, 3]. Their accuracy rests on approximately regular timing; dense mixing, harmonic ambiguity, false peaks from cross-emitter differences, and missing pulses that break sequential differencing all erode reliability [3].

Once PRI agility became widespread, many systems augmented timing with unsupervised structure in joint spaces of RF, pulse width, AOA, and related measurements, using methods such as  $k$ -means or density-based clustering [3, 4]. That shift treats deinterleaving as geometric separability in hand-crafted feature space, but performance still depends on scaling, metric choices, and hyperparameters (for example cluster counts or neighborhood radii) that are difficult to set when emitters are unknown a priori [3].

Modern emitters compound the difficulty through overlapping pulse trains, rapid mode changes, and low-observable emissions that increase loss and measurement jitter [3]. The PDW representation in Section 2.1.4 and the learning-based tools in later sections of this chapter target reasoning under those conditions.

In classical ELINT architectures, specifications therefore treat train formation as an early, largely discrete step: deinterleaving is completed first so that downstream logic receives one dominant hypothesis per physical emission path, or a short ranked list, rather than a time-interleaved mixture [1].

### 2.2.3 Emitter reporting and equipment identification

*Emitter reports* are assembled after train formation, summarizing each train with aggregated PDW statistics, tentative equipment identities from lookup tables, and temporal extent. Library-driven identification and threat correlation consume those reports together with operator context.

### 2.2.4 Operating-mode analysis and waveform libraries

Mode classification is then commonly framed as matching the train against a *mode library* of nominal RF bands, PRI/PRF families, stagger laws, and scan programs, so that declared modes remain anchored to catalogued waveform templates. When labels are incomplete, unsupervised or weakly supervised grouping over pulse trains plays the same operational role as *mode clustering* at the library boundary: it proposes coherent behavior segments that analysts later align to doctrine or intelligence products.

### 2.2.5 Emitter geolocation, fusion, and downstream displays

*Positional estimation*—for example AOA intersections from direction-finding networks, time-difference-of-arrival fixes along baselines, or fused tracks across collectors—consumes the same deinterleaved associations, which isolates geometry and timing errors from the sorting stage. Later stages may maintain libraries, threat lists, and operator displays, but for the learning problem in this thesis the critical contract is that the model consumes the deinterleaved PDW stream produced here.

## 2.3 Machine Learning Fundamentals

At its most general, ML replaces hand-coded decision rules with parametric models whose parameters are fitted to data so that a chosen performance criterion is optimized on observed examples of the task. A model is a function  $f_{\theta}: \mathcal{X} \rightarrow \mathcal{Y}$  from an input space  $\mathcal{X}$  to a task-specific output space  $\mathcal{Y}$ , and learning consists of searching over parameter vectors  $\theta$  until  $f_{\theta}$  is consistent with the regularities present in the training data. What changes between problems is the nature of the supervision available, the structure of  $\mathcal{Y}$ , and the loss that scores predictions against that supervision.

A common formal frame is *empirical risk minimization*. Given a dataset  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$  drawn from an unknown distribution and a pointwise loss  $\ell$ , the learner solves

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(f_{\theta}(\mathbf{x}_i), y_i) + \Omega(\theta), \quad (2.5)$$

where  $\Omega$  is an optional regularizer that trades fidelity to the training data against complexity of the fitted model. Generalization to fresh inputs is then assessed on held-out data, and the entire training procedure—data, model class, loss, optimizer, regularization—is judged by its effect on that out-of-sample behavior rather than on the training fit alone.

### 2.3.1 Learning paradigms

ML problems are conventionally separated by the kind of supervision available at training time. In *supervised* learning, each  $\mathbf{x}_i$  is paired with a label  $y_i$ , categorical for classification or real-valued for regression, and  $\ell$  penalizes disagreement between the prediction and that label. In *unsupervised* learning, labels are absent and the goal is to expose structure in the input distribution itself, for example through density estimation, dimensionality reduction, or clustering; Section 2.4 treats clustering specifically because it is the central downstream consumer of the embeddings used in this thesis. *Self-supervised* learning sits between the two: it constructs supervisory targets automatically from the data through a designed pretext task—predicting masked parts of an input, contrasting augmented views, or ordering shuffled tokens—so that an encoder can be trained without manual annotation. The remainder of this thesis operates in the

supervised regime: training pulse trains carry per-pulse emitter and operating-mode labels (Sections 4.1 and 4.2), and the training objective uses both label streams directly; the unsupervised and self-supervised paradigms are summarised here only because the clustering tools in Section 2.4 and the contrastive representation-learning machinery in Section 2.5 are inherited from those traditions.

### 2.3.2 Representation learning

Much of modern DL practice is best understood as *representation learning*: a parametric encoder maps high-dimensional observations into a lower-dimensional vector space whose geometry is shaped by the training objective. Downstream tasks—classification, clustering, retrieval, or sequence prediction—then operate on the resulting *embeddings*  $\mathbf{z} = f_{\theta}(\mathbf{x}) \in \mathbb{R}^d$  rather than on raw sensor coordinates. The advantage of routing tasks through a learned representation is that invariances, similarities, and task-relevant factors of variation can be encoded by  $f_{\theta}$  once, after which simple linear or geometric operations on  $\mathbf{z}$  suffice for many downstream questions. Sections 2.5 and 2.6 expand on this view; the method chapter (Chapter 4) instantiates it for PDW streams, training an encoder whose embeddings are intended to expose emitter and mode structure.

### 2.3.3 Applications in radar deinterleaving and mode classification

Radar signal deinterleaving and operating-mode classification illustrate the practical pull of these ideas. For decades, both tasks were addressed by deterministic temporal rules and hand-crafted statistics over PDW parameters, as outlined in Section 2.2. As emitters became more agile and electromagnetic environments more crowded, the field migrated first to *shallow* ML in the form of multi-parameter clustering— $k$ -means and density-based methods over RF, AOA, and pulse width—and more recently to *deep* sequence models that learn pulse-level features directly from data [3]. Recurrent and attention-based architectures now treat deinterleaving as sequence labeling in which the per-pulse output identifies an emitter, while transformer encoders in particular are well suited to capturing the non-contiguous, multi-scale relationships induced by jittered, staggered, or grouped PRIs [3, 5].

Mode classification follows a parallel trajectory. Classical ESM systems match deinterleaved pulse trains against catalogued waveform templates in a mode library [1], while learning-based approaches reframe the same task either as supervised classification over known mode labels or as unsupervised grouping over pulse trains when labels are scarce or incomplete [3]. This thesis occupies an intermediate position: training uses explicit per-pulse mode labels from a global catalogue, but the training signal is a supervised contrastive loss that shapes embedding geometry without a closed-catalogue softmax, so that previously unseen modulation types can surface as additional clusters at inference; the emitter task uses the same kind of supervised contrastive signal but with

labels whose scope is local to each pulse train, and inference also runs through clustering. The clustering tools in Section 2.4 and the contrastive representation-learning machinery in Sections 2.5 and 2.6 provide the technical foundation that the method chapter builds on.

## 2.4 Clustering Methods

Placeholder paragraph.

## 2.5 Deep Representation Learning

Many DL pipelines replace hand-crafted features with a parametric map that is trained end-to-end from data. The central object is an *encoder*  $f_{\theta}: \mathcal{X} \rightarrow \mathbb{R}^d$ , implemented in practice as a stack of nonlinear layers (for example, convolutions over a grid, recurrence over time, or attention over a sequence) followed by a pooling or readout head that yields a fixed-dimensional embedding  $\mathbf{z} = f_{\theta}(\mathbf{x})$ . The design question is which training signal makes  $\mathbf{z}$  useful for the downstream task when labels are scarce or absent.

In supervised learning,  $f_{\theta}$  is trained together with a classifier so that  $\mathbf{z}$  separates classes under a cross-entropy or margin loss. When cluster structure is the target but assignments are unknown, the same encoder can instead be trained with objectives that encourage smoothness, invariance to nuisance transformations, or agreement with a clustering hypothesis in embedding space. *Self-supervised* learning occupies an intermediate regime: the model predicts surrogate targets constructed from the input itself (for example, masked tokens, relative positions, or transformed views), so that  $f_{\theta}$  learns semantics without human labels [6].

*Contrastive* objectives implement instance discrimination by treating two augmentations of the same example as a positive pair and other examples in a minibatch (or a memory bank) as negatives. Let  $\mathbf{z}_i$  and  $\mathbf{z}_i^+$  denote normalized embeddings of two views of sample  $i$ , let  $\text{sim}(\cdot, \cdot)$  denote cosine similarity, and let  $\tau > 0$  be a temperature. A standard *InfoNCE*-style loss takes the form

$$\mathcal{L}_{\text{NCE}}(i) = -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_i^+)/\tau)}{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_i^+)/\tau) + \sum_{j \neq i} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j^+)/\tau)}, \quad (2.6)$$

which lower-bounds mutual information between representations of different views under suitable generative assumptions [7]. Large-scale implementations such as SimCLR combine strong augmentations with wide embeddings and many negatives drawn from the same minibatch [6]. Equation (2.6) should be read as a template: the positive set can be enlarged, negatives can be drawn from a queue, and the similarity can be implemented with a learned projection head without changing the underlying contrastive principle.



Pretext tasks offer an alternative route: the network solves an auxiliary problem (predicting rotations, solving jigsaw permutations, inpainting masked regions) whose solution is incidental, but the representation  $\mathbf{z}$  becomes predictive of semantically stable factors. Such objectives can be combined with clustering when the end goal is partition discovery rather than linear probe accuracy on a fixed label set. Section 3.3 surveys deep embedded clustering and prototype-based alternatives that couple encoder training with discrete structure, together with the gap relative to nested emitter–mode organization in ELINT streams.

Sequence encoders based on self-attention, including the transformer blocks reviewed in Section 2.6, instantiate  $f_{\theta}$  when  $\mathbf{x}$  is an ordered set of tokens (for example, a window of PDWs). The Method chapter builds on this stack: a contrastive or clustering-friendly loss shapes the embedding space, while the attention-based encoder models temporal context within each window.

## 2.6 The Transformer Encoder

The transformer architecture replaces recurrence with a stack of layers built around *self-attention*, so that every position in a sequence can directly attend to every other position in parallel [8]. One understanding of this is a weighted fully connected graph. Encoder-only stacks, as in BERT [9], map an input sequence to contextualised token representations of the same length and are a natural choice when the downstream task is representation learning or clustering over tokens rather than autoregressive generation.

**Scaled dot-product attention.** Attention maps three matrices of *queries*  $\mathbf{Q}$ , *keys*  $\mathbf{K}$ , and *values*  $\mathbf{V}$  to an output matrix of the same leading dimension as  $\mathbf{Q}$ . Each row of  $\mathbf{Q}$  scores all rows of  $\mathbf{K}$  through inner products; the scores are scaled, normalised with a row-wise softmax, and used to form a convex combination of the rows of  $\mathbf{V}$ . With  $d_k$  the dimension of each query and key vector (column width of  $\mathbf{Q}$  and  $\mathbf{K}$  after the usual transpose convention), the mapping is

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^{\top}}{\sqrt{d_k}}\right) \mathbf{V}. \quad (2.7)$$

The scaling by  $\sqrt{d_k}$  keeps the inner products from growing too large in magnitude as  $d_k$  increases, which would otherwise drive the softmax into near one-hot rows and yield vanishing gradients [8]. Equation (2.7) is reused later as the core attention primitive inside the encoder.

**Self-attention and MHA.** In *self-attention*,  $\mathbf{Q}$ ,  $\mathbf{K}$ , and  $\mathbf{V}$  are all linear projections of the same input sequence layout, so a token’s representation is updated by aggregating information from the full sequence. MHA runs  $h$  attention mechanisms in parallel on different learned subspaces, concatenates the head outputs, and applies one more linear

map to restore the model width [8]. The heads specialise to different relational patterns while sharing the same residual and normalisation wrapper as a single conceptual operator.

**Positional information.** Self-attention is permutation-equivariant in its inputs: reordering the rows of  $\mathbf{Q}$ ,  $\mathbf{K}$ , and  $\mathbf{V}$  in the same way reorders the output rows accordingly, but does not otherwise encode order. Transformer encoders therefore add a positional signal to token embeddings before the first layer, for example fixed sinusoidal features or learned position embeddings added element-wise to each token vector [8, 9]. When the raw inputs already carry absolute or relative timing, as is common for irregularly sampled sequences, practitioners sometimes omit a separate positional encoding and rely on those measurements instead; the method chapter returns to this choice for PDW windows.

**Encoder layer.** A standard transformer *encoder layer* alternates two sub-layers: MHA followed by a position-wise FFN applied independently at every position. Each sub-layer uses a residual connection and layer normalisation [8]. The FFN is typically a two-layer MLP with a wide inner dimension and a nonlinear activation, acting as the per-token nonlinearity while MHA handles cross-token mixing. Stacking  $N$  such layers yields a deep encoder that maps an input sequence  $(\mathbf{x}_1, \dots, \mathbf{x}_L)$  to contextualised representations  $(\mathbf{h}_1^{(N)}, \dots, \mathbf{h}_L^{(N)})$ , which downstream heads can pool or read out for sequence-level or token-level predictions.

**Computational cost.** In the usual dense implementation, each self-attention block materialises an  $L \times L$  matrix of token–token scores before the softmax, so time and activation memory for that block scale as  $\mathcal{O}(L^2)$  in the sequence length  $L$  at fixed embedding width, up to constants from the number of heads and projection ranks [8]. Repeating the block across depth multiplies that cost by the number of layers, which motivates hard caps on  $L$ , sparse or low-rank attention variants, or chunked processing when inputs are long; the method chapter states the windowing choice adopted here.

## CHAPTER 3

# Related Work

### 3.1 Classical Radar Emitter Identification

Placeholder paragraph.

### 3.2 Deep Learning for Radar Signal Analysis

Placeholder paragraph.

### 3.3 Deep Clustering and Joint Representation Learning

Classical clustering assumes that either raw features or a fixed embedding already reveal the desired partition. DEC departs from that split by learning an encoder jointly with a soft assignment to a small set of trainable cluster centroids in embedding space [10]. Training alternates between computing soft assignments from the current encoder and refining the encoder with a Kullback–Leibler objective that sharpens confident assignments toward an auxiliary target distribution, so that the network discovers both features and cluster structure. Follow-on work tightens initialization, adds reconstruction constraints, or refines the target distribution; the common theme is that discrete structure is not fixed in advance but co-adapts with the representation.

SwAV and related *prototype* methods replace dense pairwise contrastive negatives with a prediction problem onto a modest codebook of cluster prototypes, swapping predicted codes between augmented views to avoid representation collapse while still encouraging invariance [11]. They sit between purely self-supervised contrastive pre-training [6] and hard assignment  $k$ -means: gradients still flow through discrete-like targets, but the geometry is anchored by learned prototypes rather than by instance indices alone.

Several lines extend deep clustering to *multiple* partitions or hierarchical labels, for example by stacking separate assignment heads, alternating optimization over two clusterings, or enforcing orthogonality constraints between subspaces. The gap relevant to this thesis is narrower: emitter identity and operating mode are not arbitrary independent clusterings—modes nest inside emitters—yet many published pipelines still treat a single flat label space or attach two unrelated classification heads to the same backbone. Work that jointly addresses deinterleaving-style separation, in which

emitter identifiers are inherently local to each pulse train and must be recovered by clustering at inference, alongside finer-grained mode grouping that uses global mode labels for supervision but still produces clusters at inference (so that modulation types absent from the training catalogue can surface as additional clusters), all from a shared transformer encoder, remains comparatively sparse; this combination motivates the supervised dual-branch method in Chapter 4.

## **3.4 Transformers for Time Series and Sets**

Placeholder paragraph.

## **3.5 Joint and Multi-Task Clustering**

Placeholder paragraph.

## CHAPTER 4

# Method

### 4.1 Problem Formulation

#### Input

A *pulse train* is a finite, time-ordered sequence of PDWs

$$\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N), \quad \mathbf{x}_n \in \mathbb{R}^d,$$

where each  $\mathbf{x}_n$  encodes the measured parameters of one pulse (carrier frequency, amplitude, pulse width, time of arrival, and angle of arrival). The length  $N$  varies between trains.

#### Labels and induced partitions


Each pulse  $n \in \{1, \dots, N\}$  carries two labels: an emitter identity  $e_n \in \{1, \dots, K_{\text{em}}\}$  and an operating mode  $m_n \in \{1, \dots, M\}$ . The emitter symbol is unique only *within the current train*: the same integer in a different train indexes an unrelated emitter, and  $K_{\text{em}}$  varies between trains. The mode symbol is drawn from a finite *global* catalogue of size  $M$  that is shared across all trains, so  $m_n = \ell$  refers to the same mode regardless of which train the pulse comes from. However, in real-world scenarios, the actual number of distinct modes present is not fixed or known in advance, and our architecture does not require  $M$  as an input or impose a closed-set assumption.

These labels induce two partitions of the index set  $\{1, \dots, N\}$ ,

$$\mathcal{E}_k = \{n : e_n = k\}, \quad k \in \{1, \dots, K_{\text{em}}\}, \quad (4.1)$$

$$\mathcal{M}_\ell = \{n : m_n = \ell\}, \quad \ell \in \{1, \dots, M\}. \quad (4.2)$$

The sets  $\mathcal{E}_k$  form a partition of  $\{1, \dots, N\}$  with  $\mathcal{E}_k \cap \mathcal{E}_{k'} = \emptyset$  for  $k \neq k'$  and  $\bigcup_{k=1}^{K_{\text{em}}} \mathcal{E}_k = \{1, \dots, N\}$ ; every cell is non-empty by construction, because  $K_{\text{em}}$  counts only the emitters that actually appear in the current train. The sets  $\mathcal{M}_\ell$  are likewise pairwise disjoint and satisfy  $\bigcup_{\ell=1}^M \mathcal{M}_\ell = \{1, \dots, N\}$ , but their indexing runs over the entire global mode catalogue: a single train typically realises only a subset of  $\{1, \dots, M\}$ , so  $\mathcal{M}_\ell = \emptyset$  for any mode  $\ell$  absent from  $\{m_1, \dots, m_N\}$ . Mode coverage is therefore a property of a batch of trains rather than of any individual train. The first partition is defined only within a single train, while the second is comparable across trains. The two partitions are coupled: restricted to a fixed emitter  $k$ , the set  $\{m_n : n \in \mathcal{E}_k\}$



**Figure 4.1.** The two label-induced partitions of a single pulse train. Each pulse  $x_n$  carries two labels: an emitter identity  $e_n$  (encoded as fill colour) and an operating mode  $m_n$  (encoded as border style: solid, dashed, or dotted). The same index set  $\{1, \dots, N\}$  is partitioned in two distinct ways, by the emitter labelling (Equation (4.1)) and by the mode labelling (Equation (4.2)). Within each emitter subset the modes are in general mixed, and within each mode subset the emitters are in general mixed, reflecting the coupling between the two labellings.

is in general non-singleton, so  $(m_n)_{n \in \mathcal{E}_k}$  may hop between several catalogue entries. Figure 4.1 illustrates the two induced partitions on a toy pulse train of  $N = 9$  pulses, with the dual labelling rendered as a fill colour for the emitter index and a border style for the mode index.

## Dual embeddings and discrete outputs

The parametric model maps  $\mathbf{X}$  to two parallel embedding sequences,

$$\mathbf{Z}^{\text{em}} = (z_1^{\text{em}}, \dots, z_N^{\text{em}}), \quad z_n^{\text{em}} \in \mathbb{R}^p, \quad \mathbf{Z}^{\text{md}} = (z_1^{\text{md}}, \dots, z_N^{\text{md}}), \quad z_n^{\text{md}} \in \mathbb{R}^q. \quad (4.3)$$

The emitter branch  $\mathbf{Z}^{\text{em}}$  is trained so that pulses sharing an emitter are close in  $\mathbb{R}^p$  and pulses with distinct emitters are pushed apart, and the mode branch  $\mathbf{Z}^{\text{md}}$  is trained analogously against  $\{m_n\}$ . Both branches use the same batch-wide aggregation; the only asymmetry between them is the semantics of their labels, which propagates to the positive sets the loss reads (Section 4.4).

Discrete per-pulse outputs

$$\hat{e}_n \in \{1, \dots, \hat{K}_{\text{em}}\}, \quad n \in \{1, \dots, N\}, \quad (4.4)$$

$$\hat{m}_n \in \{1, \dots, \hat{M}\}, \quad n \in \{1, \dots, N\}, \quad (4.5)$$

are obtained *post hoc* by clustering  $\{z_n^{\text{em}}\}$  within each train and  $\{z_n^{\text{md}}\}$  across trains. The cardinalities  $\hat{K}_{\text{em}}$  and  $\hat{M}$  are determined by the clustering rather than fixed in advance; in particular  $\hat{M}$  may exceed  $M$  when modulation types absent from training are encountered at inference.

## Learning objective

Given a labelled collection of pulse trains, the goal is to fit the parameters of the map  $\mathbf{X} \mapsto (\mathbf{Z}^{\text{em}}, \mathbf{Z}^{\text{md}})$  by minimising a sum of two supervised metric-learning losses that share their functional form and a common batch-wide aggregation, differing only in which label drives their positive sets: emitter identifiers on  $\mathbf{Z}^{\text{em}}$  and mode labels on  $\mathbf{Z}^{\text{md}}$ . On a previously unseen train, the inference task is to recover the partitions of Equations (4.1) and (4.2) up to relabelling of the cluster symbols.

## 4.2 Data Representation and Preprocessing

Operational ELINT recordings are scarce and sensitive, so all training and evaluation data are synthetic. Pulse trains are produced by Saab’s OpenModesty scenario simulator, the same tool used by earlier KTH degree projects on radar pulse data [12]. OpenModesty specifies flight paths, emitter placements, and operating-mode schedules, then emits and intercepts pulses one at a time. Receiver-side imperfections (dropped pulses, spurious detections, parameter jitter) are produced inside the simulator and are already present in the exported pulse records, so no separate noise injection is applied downstream. Each exported pulse carries the continuous fields used here together with the two discrete labels defined in Section 4.1.

### 4.2.1 Windowing


Self-attention computes an  $L \times L$  compatibility matrix at each layer, so the computational and memory cost of a forward pass grows quadratically with the window length  $L$ , i.e.,  $\mathcal{O}(L^2)$ . Because a pulse train may contain millions of pulses, processing an entire train at once is infeasible with standard self-attention due to prohibitive memory requirements. While recent work on transformers proposes sub-quadratic attention variants to address this, the present thesis restricts itself to evaluating the baseline (plain) attention mechanism without architectural modifications.

The encoder therefore operates on fixed-length windows of  $L = 1024$  pulses. Given a pulse train  $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)$  and a stride  $s$ , the  $w$ th window is

$$\mathbf{X}^{(w)} = (\mathbf{x}_{1+(w-1)s}, \dots, \mathbf{x}_{L+(w-1)s}), \quad w \in \left\{1, \dots, \left\lfloor (N - L)/s \right\rfloor + 1\right\}. \quad (4.6)$$

Training uses  $s = L/2$ , so consecutive windows overlap by half their length; this increases the number of training views per scenario and reduces sensitivity to alignment with mode transitions. Validation and test windows use  $s = L$  so that metrics are not inflated by near-duplicate windows. The trailing remainder shorter than  $L$  is dropped. Figure 4.2 sketches the resulting segmentation.

At inference, the same fixed-length contract approximates a rolling window over a live pulse stream: a receiver buffers the most recent  $L$  pulses, applies the trained



**Figure 4.2.** Sliding-window segmentation of a pulse train as defined in Equation (4.6). The construction is valid for any stride  $0 < s \leq L$ ; the figure uses  $s = L/2$  for legibility. Successive windows  $\mathbf{X}^{(w)}$  advance by  $s$  pulses and therefore overlap by  $L - s$ . Dashed guides project window boundaries down onto the pulse train. Pulses past the right edge of the last window of length  $L$  form a trailing remainder of length  $< L$  that is dropped during training and evaluation. Training uses  $s = L/2$  and validation/test use  $s = L$ , the latter giving abutting non-overlapping windows.

encoder, advances by  $s$ , and repeats. Training-time windows are therefore representative of the operational reasoning unit, not just a memory convenience.

### 4.2.2 Per-feature scaling

Let  $\mathcal{P}_{\text{tr}}$  denote the multiset of all pulses in the training split, pooled across pulse trains. All scaling statistics are estimated on  $\mathcal{P}_{\text{tr}}$  and applied unchanged to validation and test pulses. A per-train fit would normalise out absolute parameter ranges and thereby sharpen contrasts that help intra-train deinterleaving, but operating-mode signatures are defined by those same absolute ranges and are comparable across trains by construction (Section 4.1); pooling statistics keeps the cross-train scale information that the mode branch relies on, at the price of a softer per-train normalisation for the emitter branch.

For each pulse parameter  $g \in \{f, \text{PW}, \log_{10} A\}$ , the empirical mean and standard deviation on  $\mathcal{P}_{\text{tr}}$  are

$$\mu_g = \frac{1}{|\mathcal{P}_{\text{tr}}|} \sum_{n \in \mathcal{P}_{\text{tr}}} g_n, \quad \sigma_g^2 = \frac{1}{|\mathcal{P}_{\text{tr}}|} \sum_{n \in \mathcal{P}_{\text{tr}}} (g_n - \mu_g)^2, \quad (4.7)$$

and the standardised feature is

$$\tilde{g}_n = \frac{g_n - \mu_g}{\sigma_g}. \quad (4.8)$$

Amplitude enters the pipeline as  $\log_{10} A_n$  rather than  $A_n$ . Pulse amplitudes span several orders of magnitude across emitters, ranges, and antenna patterns, and a logarithm compresses that dynamic range so that a few high-power pulses do not dominate the



standardised channel.

Time of arrival  $t_n$  is min–max scaled to  $[0, 1]$  against the training-corpus extrema,

$$\tilde{t}_n = \frac{t_n - t_{\min}}{t_{\max} - t_{\min}}, \quad (4.9)$$

with  $t_{\min}, t_{\max}$  taken over  $\mathcal{P}_{\text{tr}}$ . The azimuth and lateral components of AOA, written  $\theta_n^{\text{az}}$  and  $\theta_n^{\text{lat}}$ , are instead rescaled from their known physical intervals to  $[0, 1]$  by direct linear mapping:

$$\tilde{\theta}_n^{\text{az}} = \frac{\theta_n^{\text{az}}}{2\pi}, \quad \tilde{\theta}_n^{\text{lat}} = \frac{\theta_n^{\text{lat}}}{\pi}, \quad (4.10)$$

since  $\theta_n^{\text{az}} \in [0, 2\pi]$  and  $\theta_n^{\text{lat}} \in [0, \pi]$  by construction.

Within every window  $w$ , the scaled TOA is then offset by the TOA of the first pulse in that window,

$$\hat{t}_n = \tilde{t}_n - \tilde{t}_{1+(w-1)s}, \quad n \in \{1 + (w-1)s, \dots, L + (w-1)s\}, \quad (4.11)$$

so that the first pulse of every window has  $\hat{t} = 0$ . This removes absolute clock offset while preserving intra-window timing structure; the global min–max step fixes the unit of  $\hat{t}$ , so  $\hat{t}$  differences are comparable across windows. The AOA channels are not offset because their absolute values, not their increments, are physically meaningful.

### 4.2.3 Labels

Each pulse is assigned three labels: a train-local emitter identifier, a global operating mode label, and a pulse train ID. Only the emitter and mode labels are used during training; the pulse train ID is included solely for evaluating clustering metrics.

All labels are cast to numeric values for consistency and ease of processing.


## 4.3 Transformer Encoder Architecture

The encoder  $f_{\theta}$  takes the form of a transformer-encoder backbone [8] with a shared trunk followed by two task-specific branches; the overall layout is shown in Figure 4.3. This structure mirrors the dual-embedding formulation introduced in Section 4.1: the trunk produces a single contextualised representation of the input window, which the branches specialise into emitter-oriented and mode-oriented embeddings.

### Input embedding

Each PDW vector  $\mathbf{x}_n \in \mathbb{R}^d$ , normalised as described in Section 4.2, is mapped to a token embedding through an affine projection

$$\mathbf{h}_n^{(0)} = \mathbf{W}_{\text{in}} \mathbf{x}_n + \mathbf{b}_{\text{in}}, \quad \mathbf{h}_n^{(0)} \in \mathbb{R}^{d_{\text{model}}}, \quad (4.12)$$



**Figure 4.3.** Architecture of the encoder  $f_\theta$ . A sequence of PDWs is first lifted to  $d_{model} = 256$  by an input embedding, processed by a shared trunk of 6 transformer-encoder layers, and then split into two parallel branches with their own  $256 \rightarrow 128$  projection and 2 encoder layers. Branch outputs are mapped by linear projection heads to the emitter embedding space  $\mathbb{R}^8$  and the mode embedding space  $\mathbb{R}^{16}$ .

where  $\mathbf{W}_{\text{in}} \in \mathbb{R}^{d_{\text{model}} \times d}$  and  $d_{\text{model}}$  is the trunk width specified below. No positional encoding is added to  $\mathbf{h}_n^{(0)}$ . The temporal ordering of pulses is already carried by the TOA entry of  $\mathbf{x}_n$ , which makes a separate position signal redundant; consistent with this argument, Gunn et al. [5] report that adding sinusoidal or learnable positional encodings to a transformer deinterleaver yields a small but reproducible drop in performance, attributed to the resulting redundancy with the TOA feature.

## Shared trunk

The token sequence  $(\mathbf{h}_1^{(0)}, \dots, \mathbf{h}_L^{(0)})$  is processed by a stack of  $N_{\text{sh}}$  standard transformer-encoder layers, each comprising a MHA sub-layer and a position-wise FFN sub-layer with residual connections and layer normalisation around both. Self-attention is unmasked, so every pulse can attend to every other pulse in the window. The trunk operates with hidden dimension  $d_{\text{model}} = 256$ ,  $h_{\text{sh}} = 8$  attention heads (so that each head has dimension 32), and an inner FFN dimension of  $d_{\text{ff,sh}} = 2048$ ;  $N_{\text{sh}} = 6$  layers are used. Its output is a contextualised token sequence  $\mathbf{H}^{\text{sh}} = (\mathbf{h}_1^{\text{sh}}, \dots, \mathbf{h}_L^{\text{sh}})$  with  $\mathbf{h}_n^{\text{sh}} \in \mathbb{R}^{256}$ , in which information has been mixed across the full window.

## Task-specific branches


The trunk output feeds two parallel branches that do not share parameters, producing the dual embeddings of Equation (4.3). Each branch begins with an affine projection that reduces the hidden width from  $d_{\text{model}} = 256$  to  $d_{\text{br}} = 128$ , after which  $N_{\text{br}} = 2$  transformer-encoder layers of the same form as the trunk are applied. Within a branch, MHA uses  $h_{\text{br}} = 4$  heads (again with per-head dimension 32) and the FFN inner dimension is  $d_{\text{ff,br}} = 1024$ . The reduced width and depth reflect a deliberate asymmetry: the trunk is intended to carry the heavy lifting of contextual mixing, while the branches only need to specialise the representation to either emitter identity or operating mode. Write the resulting token sequences as  $\mathbf{H}^{\text{em}} = (\mathbf{h}_1^{\text{em}}, \dots, \mathbf{h}_L^{\text{em}})$  and  $\mathbf{H}^{\text{md}} = (\mathbf{h}_1^{\text{md}}, \dots, \mathbf{h}_L^{\text{md}})$  with  $\mathbf{h}_n^{\text{em}}, \mathbf{h}_n^{\text{md}} \in \mathbb{R}^{128}$ .

## Projection heads

Each branch terminates in a linear projection head that maps every token to its respective embedding space,

$$\mathbf{z}_n^{\text{em}} = \mathbf{W}_{\text{em}} \mathbf{h}_n^{\text{em}} + \mathbf{b}_{\text{em}}, \quad \mathbf{z}_n^{\text{md}} = \mathbf{W}_{\text{md}} \mathbf{h}_n^{\text{md}} + \mathbf{b}_{\text{md}}, \quad (4.13)$$

with  $\mathbf{z}_n^{\text{em}} \in \mathbb{R}^p$  and  $\mathbf{z}_n^{\text{md}} \in \mathbb{R}^q$  for  $p = 8$  and  $q = 16$ . The mode space is assigned the higher dimensionality because, conditional on emitter identity, mode structure is expected to require a richer geometry to separate subtle waveform regimes, whereas emitter membership induces a coarser partition over the same pulses and is well served



**Figure 4.4.** Inference-time decoding of branch embeddings into per-pulse cluster indices via Equation (4.14). Dots are pulse embeddings and dashed ellipses are DBSCAN clusters in one window. The pipeline is identical on both branches and the integer outputs are local to each clustering instance; they acquire their interpretation from the supervision scope of Section 4.1, which differs between the two branches.

by a more compact embedding.

## From continuous embeddings to discrete partitions

The projection heads in Equation (4.13) produce continuous vectors, whereas the per-pulse outputs in Equations (4.4) and (4.5) are discrete. At inference time, each branch’s  $L$ -token embedding sequence is therefore decoded independently and per window by a fixed clustering operator,

$$(\hat{e}_1, \dots, \hat{e}_L) = \text{DBSCAN}(\{z_n^{\text{em}}\}_{n=1}^L), \quad (\hat{m}_1, \dots, \hat{m}_L) = \text{DBSCAN}(\{z_n^{\text{md}}\}_{n=1}^L), \quad (4.14)$$

with cardinalities  $\hat{K}_{\text{em}}, \hat{M} \in \mathbb{N}$  inferred from the data (Figure 4.4). DBSCAN [4] groups embeddings by local density at a fixed Euclidean radius  $\varepsilon = 0.4$  and minimum neighbourhood count  $\text{minPts} = 5$ , assigning sparser points to a noise category. Clustering per window matches the train-local supervision scope of the emitter branch and keeps the runtime linear in window length; cross-window alignment of mode embeddings is instead supplied by the batch-wide contrastive term of Section 4.4 and resolved by the permutation-invariant metrics of Section 4.6. The operator contributes no gradients, and the trunk, branch, and projection-head weights together form the parameter vector  $\theta$  optimised by the loss in Section 4.4.

## 4.4 Loss Function

The training objective is a sum of two *supervised contrastive* [13] terms with the same functional form, both aggregated over the whole mini-batch. The two terms differ only in which label drives the positive set: an emitter identifier on the emitter branch, and the global mode label of Section 4.1 on the mode branch. Routing the mode branch through a contrastive loss rather than a closed-catalogue softmax keeps the inference pipeline of Equation (4.14) open to modulations absent from training; the post-hoc clustering operator contributes no gradients and is not part of the objective. And then embeddings for each PDW, or the encodings (before final projection), could probably be use as input to a know emitter library with softmax heads. This could also potentially enable transfer learning of custom classification heads from the same base model.

### Supervised contrastive kernel

All projection-head outputs of Equation (4.13) are  $\ell_2$ -normalised, so that the cosine similarity reduces to the inner product

$$s_{ij} := \mathbf{z}_i^\top \mathbf{z}_j \in [-1, 1], \quad (4.15)$$

and a fixed temperature  $\tau > 0$  rescales it inside the loss; Smaller  $\tau$  concentrates the gradient on the hardest negatives, while larger  $\tau$  stabilizes early optimization, since  $\|\nabla \mathcal{L}_{\text{SC}}\| \propto \frac{1}{\tau}$  [13]. Both branches use  $\tau = 0.2$ , the value commonly chosen for normalised contrastive features [6, 13].

For an anchor set  $I$  with embeddings  $\{\mathbf{z}_i\}_{i \in I}$ , write the candidate set  $A(i) = I \setminus \{i\}$  and the positive subset  $P(i) \subseteq A(i)$  defined by label coincidence. The supervised contrastive contribution [13] is

$$\mathcal{L}_{\text{SC}}(I, \{P(i)\}) = - \sum_{i \in I^+} \frac{1}{|P(i)|} \sum_{p \in P(i)} \log \frac{\exp(s_{ip}/\tau)}{\sum_{a \in A(i)} \exp(s_{ia}/\tau)}, \quad (4.16)$$

where  $I^+ = \{i \in I : P(i) \neq \emptyset\}$  and  $\mathcal{L}_{\text{SC}} = 0$  when  $I^+ = \emptyset$ . Anchors outside  $I^+$  contribute only as negatives. The denominator runs over every candidate in  $A(i)$ , so unrelated samples exert repulsive pressure and the encoder cannot collapse to a constant map [6].

### Branch-specific aggregation

The two branches feed Equation (4.16) with the same batch-wide index set

$$I = \{1, \dots, B\} \times \{1, \dots, L\}, \quad (4.17)$$

pooled over the  $B$  windows of a mini-batch, and differ only in the labels they consume. Let  $\mathbf{z}_{b,n}^{\text{em}}$  and  $\mathbf{z}_{b,n}^{\text{md}}$  denote the projection outputs at position  $n \in \{1, \dots, L\}$  in window  $b \in \{1, \dots, B\}$ , with emitter label  $e_{b,n}$  and mode label  $m_{b,n}$ . Emitter identifiers are train-local in Section 4.1 and are taken to be globally unique within a mini-batch by pairing the train-local integer with the originating pulse-train identifier, so that  $e_{b',n'} = e_{b,n}$  iff the two pulses come from the same physical emitter. The two positive sets are

$$P_{b,n}^{\text{em}} = \{(b', n') \in I \setminus \{(b, n)\} : e_{b',n'} = e_{b,n}\}, \quad P_{b,n}^{\text{md}} = \{(b', n') \in I \setminus \{(b, n)\} : m_{b',n'} = m_{b,n}\}, \quad (4.18)$$

and the two branch losses are

$$\mathcal{L}_{\text{em}} = \mathcal{L}_{\text{SC}}(I, \{P_{b,n}^{\text{em}}\}), \quad (4.19)$$

$$\mathcal{L}_{\text{md}} = \mathcal{L}_{\text{SC}}(I, \{P_{b,n}^{\text{md}}\}). \quad (4.20)$$

Both losses are therefore standard SupCon [13] aggregated over the whole batch. Because the loss only consumes pairwise label coincidences, neither the emitter count  $K_{\text{em}}$  nor the catalogue size  $M$  appears in Equation (4.16), and modes held out at training time still inherit useful geometry from their neighbours in  $\mathbb{R}^q$ . The total training objective is

$$\mathcal{L} = \mathcal{L}_{\text{em}} + \lambda_{\text{md}} \mathcal{L}_{\text{md}}, \quad \lambda_{\text{md}} > 0, \quad (4.21)$$

with  $\lambda_{\text{md}}$  reported in Chapter 5.

## Probabilistic interpretation

The kernel Equation (4.16) admits a Gibbs–Boltzmann reading that clarifies the role of every factor. For each anchor  $i \in I^+$ , define the categorical distribution over its candidates

$$P_{\theta}(a \mid i) = \frac{\exp(s_{ia}/\tau)}{Z_i}, \quad Z_i = \sum_{a' \in A(i)} \exp(s_{ia'}/\tau), \quad a \in A(i), \quad (4.22)$$

which is exactly the Boltzmann distribution with energy  $-s_{ia}$  and temperature  $\tau$  on the configuration space  $A(i)$ . Substituting Equation (4.22) into Equation (4.16) yields the per-anchor cross-entropy between the empirical positive distribution (uniform on  $P(i)$ ) and the model  $P_{\theta}(\cdot \mid i)$ ,

$$\mathcal{L}_{\text{SC}} = - \sum_{i \in I^+} \frac{1}{|P(i)|} \sum_{p \in P(i)} \log P_{\theta}(p \mid i), \quad (4.23)$$

that is, minus the average log-evidence the encoder assigns to the positives, exactly as a maximum-likelihood estimator would treat  $P(i)$  as a sample of observations drawn from the latent positive class.

Applying  $\log(x/y) = \log x - \log y$  to Equation (4.16) separates the attractive and repulsive contributions explicitly:

$$\mathcal{L}_{\text{SC}} = \underbrace{\sum_{i \in I^+} \log Z_i}_{\text{log-partition (repulsion)}} - \underbrace{\sum_{i \in I^+} \frac{1}{|P(i)|} \sum_{p \in P(i)} \frac{s_{ip}}{\tau}}_{\text{mean positive similarity (attraction)}}. \quad (4.24)$$

The first term is monotone in similarities to every candidate, with  $\partial \log Z_i / \partial s_{ia} = P_{\theta}(a | i) / \tau$ , so the gradient automatically concentrates on the hardest negatives—those whose Boltzmann weight under Equation (4.22) is largest despite carrying a different label. The second term rewards alignment of each anchor with its positives, with each positive contributing  $1/(\tau|P(i)|)$  to the gradient.

The prefactor  $1/|P(i)|$  turns the inner sum of log-probabilities into an arithmetic mean, equivalently into the log of a geometric mean,

$$-\frac{1}{|P(i)|} \sum_{p \in P(i)} \log P_{\theta}(p | i) = -\log \left( \prod_{p \in P(i)} P_{\theta}(p | i) \right)^{1/|P(i)|}. \quad (4.25)$$

Without the prefactor, the inner sum equals the log-likelihood of  $|P(i)|$  conditionally independent observations of the latent positive class, and the loss would scale extensively with the number of available positives. Dividing by  $|P(i)|$  replaces this product by its geometric mean, so an anchor’s contribution becomes intensive in the size of its positive set; this is what allows the same kernel to remain well-behaved across the two branches, where  $|P(i)|$  ranges from a handful (emitter contrasts inside one window) to many hundreds (mode contrasts pooled over the batch).

## 4.5 Training Procedure

The parameter vector  $\theta$  attached to the maps in Equations (4.12) and (4.13) is estimated by minimising the training objective  $\mathcal{L}$  in Equation (4.21) on mini-batches of windows. Gradients are obtained by reverse-mode automatic differentiation and used in an adaptive first-order optimiser of the Adam family [14]. Unless Section 5.3 specifies otherwise, the implementation follows the default moment decay rates  $(\beta_1, \beta_2) = (0.9, 0.999)$ , stabiliser  $\epsilon = 10^{-8}$ , and bias-corrected running estimates of the first and second moments described by Kingma and Ba [14], as exposed by the PyTorch optimiser interface.

The learning rate follows a cosine decay from a peak value down to a small floor over the course of training, which keeps step sizes comparatively large while the loss landscape is coarse and shrinks them during the final refinement phase. Peak learning rate, weight decay, batch size, number of passes over the training windows, and the exact step counting convention (per epoch versus global step) are hyperparameters and are listed in Table 5.2. At the start of each epoch, training windows are shuffled so that consecutive mini-batches are not ordered by scenario; any fixed random seeds,

data-loader worker counts, and reproducibility settings used in the reported runs are stated in Sections 5.1 and 5.3.

Dropout inside the transformer blocks of Section 4.3 follows the usual pattern of stochastic attention and feed-forward masks; dropout rates belong to the same hyperparameter table. Training can be stopped early when a validation score computed on held-out windows (for example the smoothed mini-batch loss or an auxiliary clustering criterion) fails to improve for a fixed patience measured in epochs, in which case the best checkpoint prior to stagnation is retained; patience, validation metric, and checkpoint policy are also recorded in Table 5.2.

The mode-branch weight  $\lambda_{\text{md}}$  may be held small during an initial warm-up segment so that the trunk first learns coarse emitter geometry from the within-train supervised contrastive term in Section 4.4 before the batch-wide mode-contrastive gradients dominate; whether such a ramp is used, and its length, is part of the same configuration table. Multi-phase schedules, such as representation pretraining followed by joint fine-tuning, are compatible with the same tooling; if they are employed, the phases and their objectives are stated in Section 5.3.

## 4.6 Evaluation Metrics

The methods of Sections 4.3 and 4.4 return per-pulse cluster indices  $\hat{y}_n$  that have to be compared with reference labels  $y_n$ . Both labellings partition the same index set  $\{1, \dots, N\}$ , instantiated by either the emitter pair  $(e_n, \hat{e}_n)$  within a window or the mode pair  $(m_n, \hat{m}_n)$  across windows, as in Equations (4.1) to (4.3). The integers used to name each cluster are arbitrary and may differ between the two labellings, so every metric reported here is invariant to relabelling: only the induced partitions matter, not the symbols themselves. This is essential because, by construction, the clustering operator of Equation (4.14) produces window-local integers that bear no fixed relation to the ground-truth catalogue.

All metrics derive from the contingency table

$$n_{ij} = \sum_{n=1}^N \mathbf{1}[y_n = i, \hat{y}_n = j],$$

with marginals  $a_i = \sum_j n_{ij}$  and  $b_j = \sum_i n_{ij}$  counting how many pulses carry true label  $i$  or predicted label  $j$ . No single scalar captures every facet of cluster agreement, so we report metrics drawn from three different families — pair counting, information theory, and entropy-based homogeneity/completeness — together with a Hungarian-aligned accuracy when the two cardinalities are comparable.



## Pair-counting view: Rand and adjusted Rand

The most intuitive way to compare two partitions is to look at every *pair* of pulses and ask whether the two labellings agree on a binary question: are these two pulses in the *same* cluster, or in *different* clusters? Each unordered pair  $\{m, n\}$  falls into one of four categories: same in both, same in  $Y$  but different in  $\hat{Y}$ , different in  $Y$  but same in  $\hat{Y}$ , or different in both. The first and last are agreements; the middle two are disagreements. A perfectly aligned pair of partitions has only agreements; a random pair has many disagreements.

It is convenient to count pairs that are placed *together* by each partition. Writing  $S_{Y\hat{Y}}$ ,  $S_Y$ , and  $S_{\hat{Y}}$  for the numbers of pairs that are together in both partitions, in  $Y$ , and in  $\hat{Y}$  respectively, the contingency table gives

$$S_{Y\hat{Y}} = \sum_{i,j} \binom{n_{ij}}{2}, \quad S_Y = \sum_i \binom{a_i}{2}, \quad S_{\hat{Y}} = \sum_j \binom{b_j}{2}, \quad (4.26)$$

out of  $\binom{N}{2}$  total pairs. The plain Rand index [15] is the fraction of pairs on which the two partitions agree,

$$R(Y, \hat{Y}) = 1 - \frac{S_Y + S_{\hat{Y}} - 2S_{Y\hat{Y}}}{\binom{N}{2}} \in [0, 1]. \quad (4.27)$$

The Rand index is easy to read but inflated: when most pulses sit in different clusters, the bulk of pairs are correctly judged “different in both” simply by chance, so  $R$  stays close to one even for poorly aligned partitions. The ARI [15] corrects this by subtracting the expected value of  $S_{Y\hat{Y}}$  under a permutation null that fixes the cluster sizes  $\{a_i\}$  and  $\{b_j\}$  but otherwise reassigns pulses uniformly,

$$\mathbb{E}[S_{Y\hat{Y}}] = \frac{S_Y S_{\hat{Y}}}{\binom{N}{2}}. \quad (4.28)$$

Subtracting this expectation from numerator and denominator and rescaling so that perfect agreement gives one yields

$$\text{ARI} = \frac{S_{Y\hat{Y}} - \mathbb{E}[S_{Y\hat{Y}}]}{\frac{1}{2}(S_Y + S_{\hat{Y}}) - \mathbb{E}[S_{Y\hat{Y}}]}, \quad \text{ARI} \leq 1. \quad (4.29)$$

The numerator is the excess of together-together pairs over chance; the denominator is the largest such excess achievable while preserving the marginals. Hence  $\text{ARI} = 1$  at exact recovery up to relabelling,  $\text{ARI} \approx 0$  at chance, and negative values indicate worse-than-random alignment.

## Information-theoretic view: NMI and AMI

A second viewpoint asks how much knowing one labelling tells us about the other: if  $\hat{Y}$  is a deterministic function of  $Y$ , then announcing  $\hat{Y}$  leaves no remaining uncertainty about  $Y$ , and conversely two independent labellings are uninformative about each other. This is exactly what mutual information measures [16]. With Shannon entropies of the marginals

$$H(Y) = - \sum_{i: a_i > 0} \frac{a_i}{N} \log \frac{a_i}{N}, \quad H(\hat{Y}) = - \sum_{j: b_j > 0} \frac{b_j}{N} \log \frac{b_j}{N}, \quad (4.30)$$

quantifying the marginal uncertainty of each labelling, the empirical mutual information is

$$\text{MI}(Y; \hat{Y}) = \sum_{i,j: n_{ij} > 0} \frac{n_{ij}}{N} \log \frac{N n_{ij}}{a_i b_j}, \quad (4.31)$$

with the convention  $0 \log 0 = 0$ . MI is non-negative, zero when  $Y$  and  $\hat{Y}$  are independent, and equal to either marginal entropy when one labelling determines the other.

The NMI bounds MI to  $[0, 1]$  by dividing through by a normaliser built from the marginal entropies; we use the arithmetic mean [16],

$$\text{NMI}(Y; \hat{Y}) = \frac{2 \text{MI}(Y; \hat{Y})}{H(Y) + H(\hat{Y})}, \quad H(Y) + H(\hat{Y}) > 0, \quad (4.32)$$

and set  $\text{NMI} = 0$  if both entropies vanish. Like the Rand index, MI has a non-zero baseline expectation, and that expectation grows with the number of clusters; the AMI [16] performs the analogous chance correction. Let  $\mathbb{E}[\text{MI}]$  be the expectation of MI under a permutation null with fixed cluster sizes; then

$$\text{AMI}(Y; \hat{Y}) = \frac{\text{MI}(Y; \hat{Y}) - \mathbb{E}[\text{MI}]}{\frac{1}{2}(H(Y) + H(\hat{Y})) - \mathbb{E}[\text{MI}]}, \quad (4.33)$$

when the denominator is positive; otherwise  $\text{AMI} \equiv 0$ . AMI inherits the same range and chance baseline as ARI but operates on the information-theoretic axis, which makes it less sensitive to differences in cluster size and more sensitive to fine-grained mislabellings.

## Homogeneity, completeness, and V-measure

A third viewpoint asks two complementary questions about cluster purity: *homogeneity*, “does each predicted cluster contain only one true class?”, and *completeness*, “is each true class collected into a single predicted cluster?”. Splitting a true class across many predicted clusters preserves homogeneity but destroys completeness; merging several true classes into one predicted cluster preserves completeness but destroys homogeneity. Reporting both quantities separately therefore diagnoses whether the encoder oversplits or undermerges.

With the conditional entropies

$$H(Y | \hat{Y}) = - \sum_{i,j: n_{ij} > 0} \frac{n_{ij}}{N} \log \frac{n_{ij}}{b_j}, \quad H(\hat{Y} | Y) = - \sum_{i,j: n_{ij} > 0} \frac{n_{ij}}{N} \log \frac{n_{ij}}{a_i}, \quad (4.34)$$

homogeneity and completeness are [17]

$$h = 1 - \frac{H(Y | \hat{Y})}{H(Y)}, \quad c = 1 - \frac{H(\hat{Y} | Y)}{H(\hat{Y})}, \quad (4.35)$$

with  $h = 1$  if  $H(Y) = 0$  and  $c = 1$  if  $H(\hat{Y}) = 0$ . The V-measure is the harmonic mean

$$V(Y, \hat{Y}) = \frac{2hc}{h+c}, \quad h+c > 0, \quad (4.36)$$

and  $V = 0$  if  $h = c = 0$ . The harmonic mean is small whenever either component is small, so a high V-measure requires both homogeneity and completeness to be high simultaneously.

## Clustering accuracy with Hungarian alignment

The most directly interpretable scalar is the fraction of pulses correctly classified, but the cluster integers are arbitrary so a one-to-one relabelling of  $\hat{Y}$  onto  $Y$  must be chosen first. Among all such relabellings, the one that maximises the number of correct assignments is the optimal solution of a linear assignment problem on the contingency table  $\{n_{ij}\}$ , which the Hungarian algorithm [18] solves in polynomial time. The ACC is the resulting fraction,

$$\text{ACC}(Y, \hat{Y}) = \frac{1}{N} \max_{\pi} \sum_{n=1}^N \mathbf{1}[y_n = \pi(\hat{y}_n)], \quad (4.37)$$

where the maximum is taken over injective maps  $\pi$  between predicted and true cluster identifiers; when the two cardinalities differ, the contingency table is zero-padded so that  $\pi$  remains injective.  $\text{ACC} \in [0, 1]$ , with  $\text{ACC} = 1$  at perfect recovery up to relabelling. Compared with the other metrics, ACC is the most familiar quantity but the least forgiving when  $\hat{K}$  deviates strongly from the true number of clusters, because unmatched predicted clusters then contribute only zeros to the maximum.

## What each metric does and does not tell us

The four metrics agree at the extremes — all return one at exact recovery and approximately zero at chance — but they disagree in the middle. ARI summarises the pair structure and is comparable across runs with different cluster counts. AMI and NMI capture how much information one labelling carries about the other; AMI is preferable when the number of clusters varies, while NMI remains a useful reference because it is

universally reported. Reporting homogeneity and completeness alongside the V-measure separates oversplitting from undermerging, which matters for the open-set mode setting where  $\hat{M} > M$  is expected. ACC is reported when  $\hat{K}$  and the true number of clusters are comparable; it is the easiest number to communicate but the most fragile when cardinalities diverge. We therefore report several scores rather than a single headline metric.

## Qualitative inspection

Two-dimensional projections of  $\{z_n^{\text{em}}\}$  and  $\{z_n^{\text{md}}\}$  coloured by  $(e_n, m_n)$  complement the quantitative scores of Equations (4.29), (4.32), (4.33), (4.36) and (4.37); protocols and numbers appear in Chapter 5.

## 4.7 Baseline Methods

The baselines isolate the contribution of the dual-branch layout in Section 4.3 and of the mode term  $\lambda_{\text{md}}\mathcal{L}_{\text{md}}$  in Equation (4.21) relative to a single-stack encoder and to classical clustering without a sequence model. Final configurations and the full list of comparisons are deferred to Chapter 5; this section fixes only the design intent of each.

**Single-stack transformer.** A flat encoder of  $N_{\text{ss}}=8$  transformer-encoder layers with the same width, heads, and FFN expansion as the trunk in Section 4.3, terminated by one linear head  $\mathbb{R}^{d_{\text{model}}} \rightarrow \mathbb{R}^p$  and trained on  $\mathcal{L}_{\text{em}}$  alone. Tests whether a single representation optimised purely for window-level deinterleaving also serves the mode clustering.

**HDBSCAN on fixed window features.** No transformer; HDBSCAN [19] clusters a deterministic vector built from the scaled PDWs of a window (vectorisation specified in Chapter 5). Separates the gain from learned sequence context from what density clustering already recovers on hand-specified summaries.

**External references.**  $k$ -means on raw window vectors and on autoencoder embeddings, DEC-style objectives, SwAV [11] assignment prediction, and a closed-catalogue supervised classifier as an upper bound on the modes seen during training. The supervised classifier is reported only as a reference ceiling; it cannot accommodate the held-out modulations of the open-set experiments and is not used as the fielded mode predictor.

## CHAPTER 5

# Experiments and Results

### 5.1 Experimental Setup

All neural models in this chapter were trained and evaluated on a single workstation with one NVIDIA A100 GPU (80 GB high-bandwidth memory; some system utilities report a capacity closer to 82 GB because of binary versus decimal gigabyte conventions). The implementation uses Python with PyTorch; exact package versions are pinned in the environment file shipped with the thesis code repository when that repository is published.

Random seeds control weight initialization, training-window shuffling each epoch, and stochastic augmentations. Unless an experiment explicitly varies the seed, the same seed set is reused so that ablations in Section 5.5 are comparable under identical noise conditions. Concrete seed integers, per-run overrides, and mean  $\pm$  standard deviation over multiple seeds for quantitative scores are reported with the tables in Section 5.4.

### 5.2 Datasets

All experiments in this chapter use synthetic PDW streams generated by the proprietary scenario simulator introduced in Section 4.2. The simulator supplies time-ordered pulse records together with ground-truth labels for emitter identity (train-local, in the sense of Section 4.1) and operating mode (drawn from the global catalogue); both label streams are consumed by the training objective in Section 4.4 and are also used to compute the evaluation metrics in Sections 4.6 and 5.4. Receiver-side degradations (dropped pulses, spurious detections, parameter jitter) are produced inside the simulator rather than injected afterwards (Section 4.2); the tables below state which simulator regime is active for each run. Operational ELINT recordings are not included, which sidesteps distribution shift and disclosure constraints while still exercising mixture traffic, mode changes within emitters, and variable train density.

Each *scenario* is one simulated timeline: a sequence of PDWs whose length and the number of concurrent emitters vary across draws so that the encoder sees both sparse and crowded electromagnetic pictures. Scenarios are disjointly assigned to training, validation, and test partitions before any windowing; all normalisation statistics in Section 4.2 are fit on the training partition only and applied unchanged to validation and test pulses. Training windows use stride  $L/2$  with  $L = 1024$  pulses, whereas validation and test windows use stride  $L$ , so the effective number of windows per scenario differs by split even when raw pulse counts are similar.

**Table 5.1.** Synthetic scenario corpus used for training and evaluation. Dashes mark quantities still being reconciled with the export manifest.

| Split      | Scenarios | Pulses (approx.) | Distinct emitters | Distinct modes |
|------------|-----------|------------------|-------------------|----------------|
| Train      | —         | —                | —                 | —              |
| Validation | —         | —                | —                 | —              |
| Test       | —         | —                | —                 | —              |
| All        | —         | —                | —                 | —              |

Table 5.1 summarises corpus-level figures. The mode column counts distinct global mode symbols that appear anywhere in the corresponding split; it represents the training-time catalogue size  $M$  in Section 4.1 for the training split, while validation and test splits may exercise additional modulation variants when held-out-mode experiments are active in Chapter 5. The emitter column counts the union of train-local emitter symbols, summed across scenarios in the split: because emitter identifiers are unique only within a pulse train (Section 4.1), this number describes total emitter incidence rather than a global set, and the within-window clustering at inference operates on the train-local sub-counts. Scenario design and traffic rules also skew how often each emitter and each mode is exercised, so aggregate pulse counts need not match intuitive “balanced” priors even when the simulator exposes many labels. Because the encoder sees only windows of fixed length  $L$  (Section 4.2), rare combinations of emitter, mode, and interleaving pattern can be missing from individual windows or represented by only a handful of pulses, which limits what any single forward pass can evidence about those labels. The numeric entries are placeholders until the final export from the simulator pipeline is frozen; the intended reporting convention is scenarios for volume, pulses for aggregate exposure, and distinct labels for emitter and mode coverage.

### 5.2.1 Qualitative properties of the synthetic corpus

Beyond aggregate counts, the corpus is checked qualitatively before large training runs: example timelines are inspected for plausible PRI structure per mode, for physically admissible RF and pulse-width ranges, and for the intended mixture density when multiple emitters are active. When the simulator regime adds dropped or spurious pulses (Section 4.2), spot checks confirm that surviving pulses retain their original time order and that spurious insertions stay within the configured parameter ranges. This mirrors the reporting practice in earlier Saab-hosted simulator studies, which often include illustrative emitter tracks alongside tables [12]. Section 5.6 complements the brief sanity checks here with embedding-space visualizations and other qualitative diagnostics once model checkpoints are available.

**Table 5.2.** Primary optimisation hyperparameters for the proposed model.

| Setting                                | Value                                                                                             |
|----------------------------------------|---------------------------------------------------------------------------------------------------|
| Optimiser                              | Adam [14] with PyTorch defaults<br>( $\beta_1, \beta_2$ ) = (0.9, 0.999) and $\epsilon = 10^{-8}$ |
| Peak learning rate $\eta_{\max}$       | —                                                                                                 |
| Minimum learning rate $\eta_{\min}$    | 0 (cosine floor; PyTorch default)                                                                 |
| Cosine period $T_{\max}$               | 30 scheduler steps, one step per training epoch                                                   |
| Maximum epochs                         | 30                                                                                                |
| Early stopping patience                | 7 epochs without improvement on the validation loss                                               |
| Early stopping checkpoint              | Weights from the epoch with lowest validation loss so far                                         |
| Dropout (attention and feed-forward)   | 0.1 throughout the transformer blocks                                                             |
| Batch size $B$ (windows per step)      | —                                                                                                 |
| Mode loss weight $\lambda_{\text{md}}$ | —                                                                                                 |
| Weight decay                           | —                                                                                                 |

### 5.3 Hyperparameters and Training Details

This section collects numerical training choices deferred from Sections 4.3 to 4.5. Unless Section 5.5 states otherwise, baseline models use the same epoch budget, hardware (Section 5.1), and early-stopping protocol so that differences reflect architecture and objective rather than compute.

The contrastive temperature is fixed at  $\tau = 0.2$  as in Section 4.4. Table 5.2 summarises the optimisation schedule and regularisation for the proposed dual-branch encoder; entries marked with an em dash are still being matched to the released training script and will be filled before the final thesis submission.

The cosine scheduler follows the PyTorch `CosineAnnealingLR` convention: after each epoch the learning rate is updated from  $\eta_{\max}$  toward  $\eta_{\min}$  over  $T_{\max} = 30$  steps, which aligns the period with the nominal training horizon. If early stopping triggers first, training halts at the best validation checkpoint and the final few scheduler steps may not be executed.

### 5.4 Quantitative Results

This section reports primary clustering scores for the proposed model and for each baseline under the data conditions in Section 5.2. Rows index methods (including ablated variants where applicable); columns index emitter-level and mode-level NMI, ARI, and ACC after optimal alignment of cluster indices to the reference simulator labels (Section 4.6). Each entry is aggregated over random seeds and window splits as

stated in Sections 5.1 and 5.3; mean  $\pm$  one standard deviation is shown when multiple seeds are available.

## 5.5 Ablation Studies

Ablations isolate the contribution of individual components in Chapter 4 by retraining with exactly one change relative to the full configuration (matched data, seeds, and hyperparameters unless the ablation itself is a hyperparameter). Planned comparisons include removing or downweighting the emitter contrastive term, the mode contrastive term, or the coupling between branches; reducing the number of attention layers or heads; disabling stochastic augmentations beyond deterministic scaling; and swapping HDBSCAN post-processing for a simpler fixed- $k$  assignment when that variant is implemented.

## 5.6 Qualitative Analysis

Qualitative analysis complements the scalar metrics in Section 5.4 by checking whether embeddings organize pulses in ways that align with physical emitter boundaries and with mode transitions inside a window. The checks in Section 5.2.1 validate that the simulator export behaves as intended; this section applies analogous scrutiny *after* models are trained.

Two-dimensional projections of pooled or per-pulse embeddings (for example UMAP or t-SNE) should be colored separately by reference emitter (using the train-local emitter label of each pulse) and by reference mode (using the global mode catalogue from Section 4.1). Disjoint emitter-colored manifolds with mode-colored substructure that stays coherent within an emitter are evidence that the dual geometry in Chapter 4 is behaving as designed; elongated bridges between emitter manifolds often indicate residual deinterleaving ambiguity or shared waveform parameters across platforms.

Attention maps or head-wise statistics, if exported, can be overlaid on TOA-ordered pulses to see whether specific heads track stable intervals versus rapid mode hops.

Confusion matrices between inferred clusters and reference labels (after optimal assignment) are optional but useful when a small number of emitter or mode classes dominate error mass.



## CHAPTER 6

# Discussion

### 6.1 Interpretation of Results

This section connects the quantitative tables in Section 5.4 and the qualitative material in Section 5.6 to the modeling choices in Chapter 4. The goal is not to restate individual scores, but to explain *which* mechanisms plausibly drive the observed ranking of methods and ablations.

When emitter-level NMI or ARI lag while mode-level scores are stronger, a natural hypothesis is that the encoder prioritizes intra-window structure tied to waveform parameters (which modes share) before it fully exploits slower cues that separate physical platforms across mixed traffic. Conversely, if emitter scores dominate but mode partitions collapse, the dual-head coupling or the mode contrastive term may be underweighted relative to the emitter branch, or certain modes may be too short-lived inside length- $L$  windows to support stable clusters.

The dropped and spurious-pulse regimes produced by the simulator (Section 4.2) interact with self-attention in two ways: they perturb local timing patterns that define modes, and they change which pulses co-occur inside a window, which can confuse any window-level emitter summary. Interpretation should therefore read noise sweeps alongside the corpus diagnostics in Sections 5.2.1 and 5.6: a metric drop that coincides with visibly corrupted tracks is evidence of a robustness limit, whereas a drop on clean held-out scenarios points toward generalization or optimization issues instead.

### 6.2 Comparison with Related Work

Relative to classical ELINT fingerprinter pipelines surveyed in Section 3.1, the proposed approach trades hand-engineered feature logic for a learned encoder, at the cost of data hunger, compute, and reduced transparency unless accompanied by diagnostics such as those in Section 5.6. Unlike the supervised radar classifiers discussed in Section 3.2, which typically assume a fixed global label set for emitters and a fixed mode catalogue at inference, the training objective here treats emitter identity as a *train-local* symbol used only through within-train comparisons (Section 4.4); operating-mode labels are used globally during training to shape the mode embedding geometry, but the inference path runs through clustering rather than through a softmax over the training catalogue, so that the same model can accommodate modulation types absent from the training data.

Deep clustering methods in Section 3.3 typically optimise a *single* flat partition

without labels, or attach two unrelated supervised heads to the same backbone. This thesis instead targets two aligned partitions with a nesting interpretation (modes within emitters), with two supervised contrastive losses that share the same functional form and differ only in the scope of their positive sets: within-window for emitter labels and batch-wide for global mode labels. The setup is closer in spirit to offline pulse-level grouping work that reasons about geometry and overlap in ESM data [20] but uses gradient-based representation learning rather than hand-tuned density stages alone. Where scores fall short of the best published supervised benchmarks, the gap is expected: those models often operate on curated single-emitter tracks with a flat, globally consistent emitter label set, whereas the present setting stresses mixture windows and a label structure that forces emitter inference through clustering rather than through a global classifier.

### 6.3 Limitations

The experimental programme is offline: models are trained and scored on bounded, revisit-able windows rather than under the latency, memory, and non-stationarity constraints of a continuously arriving pulse stream (Sections 1.4, 1.4.2 and 4.2). That setting supports controlled optimization and reproducible metrics but does not demonstrate that the same representations would support stable online emitter or mode tracking without periodic retraining or revised windowing policy.

Fixed-length windows are partly a consequence of full self-attention, whose standard implementation scales quadratically with the number of pulses in the window (Section 4.2). The chosen cap on  $L$  therefore trades faithful coverage of long mode segments and rare emitters against compute, and the present pipeline drops trailing segments shorter than  $L$ , which can remove evidence entirely when a scenario does not fill an integral number of windows.

Interleaved emitters and mode-hopping further skew how many pulses from each ground-truth label appear inside any given window, so performance summaries aggregate over heterogeneous label incidence rather than over an idealized balanced design (Section 5.2). Where metrics depend on comparing inferred partitions to reference labels, conclusions are most informative for label configurations that actually recur with sufficient mass inside windows; weak support for a label in a window is a limitation of the observation model, not necessarily of the encoder in isolation.

Finally, all quantitative evidence is derived from synthetic scenarios with full reference labels; operational recordings with imperfect deinterleaving, sensor dropouts, and distribution shift are not included here. Related Saab-line work on track association under deinterleaving errors highlights that real pipelines fragment pulse trains and introduce structured mistakes that are only partly captured by the stylized drop and spurious models used for stress tests [21]. Likewise, classifiers trained on simulator draws can exhibit overconfidence on held-out emitters or modulation families, which motivates explicit out-of-distribution and drift analyses when moving toward fielded

receivers [22]. The present thesis does not implement drift monitors or uncertainty heads; those remain orthogonal safeguards if embeddings from Chapter 4 are ever deployed under evolving threat libraries.

The HDBSCAN post-processing used in Section 4.3 also encodes density assumptions that may not match every deployment, and we do not exhaust efficient-attention or hierarchical alternatives that could extend context within the same hardware envelope.

## 6.4 Implications

For ELINT and ESM practice, the main takeaway is methodological: joint emitter- and mode-aware embeddings can be learned from PDW streams whose supervision is structured asymmetrically (train-local emitter labels, global mode catalogue), with downstream discretisation handled by clustering on *both* branches so that previously unseen emitters and modulation types can in principle surface as distinct clusters rather than be silently absorbed into existing catalogue entries. Operators should treat the resulting per-pulse emitter and mode cluster indices as hypotheses that require human or library-based adjudication before tactical use, especially when training data are synthetic and drift relative to live emitters or to mode catalogues is unmodelled; in particular, mode-branch clusters whose count exceeds the size of the training catalogue, or that occupy regions of the mode embedding space far from those populated by training-time modes, are informative as candidate novel modulations and warrant separate review.

For the deep-clustering and sequence-modeling research communities, the thesis sits at the intersection of contrastive representation learning on irregular pulse sequences and multi-partition clustering with a physical nesting constraint (modes within emitters). The architectural and loss choices in Chapter 4 are portable to other domains with hierarchical latent structure, provided that evaluation metrics separate which level of the hierarchy is being scored.

Follow-up work naturally includes streaming or sliding-window evaluation, richer deinterleaving and track-error models [21], and explicit ML safeguards against dataset shift and overconfident extrapolation [22]. On the ethics side, the considerations in Section 1.6 still apply: dual-use sensitivity, data stewardship if operational traces are introduced later, and the energy cost of transformer-scale training should be revisited whenever the experimental footprint grows.

## **CHAPTER 7**

# **Conclusion**

### **7.1 Summary**

Placeholder paragraph.

### **7.2 Answers to the Research Questions**

Placeholder paragraph.

### **7.3 Future Work**

Placeholder paragraph.

# References

- [1] R. G. Wiley, *ELINT: The Interception and Analysis of Radar Signals*. Artech House, 2006.
- [2] M. I. Skolnik, *Radar Handbook*, 3rd ed. McGraw-Hill, 2008.
- [3] Z. Qu, J. Zhang, Y. Zhou, and L. Ni, “The intelligent evolution of radar signal deinterleaving: A systematic review from foundational algorithms to cognitive AI frontiers”, *Sensors*, vol. 26, no. 1, 2026. DOI: 10.3390/s26010248
- [4] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu, “A density-based algorithm for discovering clusters in large spatial databases with noise”, in *Proceedings of the 2nd International Conference on Knowledge Discovery and Data Mining (KDD)*, 1996, pp. 226–231.
- [5] E. Gunn, A. Hosford, D. Mannion, J. Williams, V. Chhabra, and V. Nockles, *Radar pulse deinterleaving with transformer based deep metric learning*, arXiv preprint arXiv:2503.13476, 2025. arXiv: 2503.13476 [cs.LG].
- [6] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, “A simple framework for contrastive learning of visual representations”, in *International Conference on Machine Learning (ICML)*, 2020.
- [7] A. van den Oord, Y. Li, and O. Vinyals, *Representation learning with contrastive predictive coding*, arXiv preprint arXiv:1807.03748, 2018. arXiv: 1807.03748 [cs.LG].
- [8] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 5998–6008.
- [9] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding”, in *Proceedings of NAACL-HLT*, 2019, pp. 4171–4186.
- [10] J. Xie, R. Girshick, and A. Farhadi, “Unsupervised deep embedding for clustering analysis”, in *International Conference on Machine Learning (ICML)*, 2016, pp. 478–487.
- [11] M. Caron, I. Misra, J. Mairal, P. Goyal, P. Bojanowski, and A. Joulin, “Unsupervised learning of visual features by contrasting cluster assignments”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [12] T. Jonsson, “Classification of Radar Emitters using Semi-Supervised Contrastive Learning”, Host company Saab AB; supervised learning experiments on pulse data from the OpenModesty scenario simulator, Master’s thesis, KTH Royal Institute of Technology, Stockholm, Sweden, 2023.

- [13] P. Khosla, P. Teterwak, C. Wang, A. Sarna, Y. Tian, P. Isola, A. Maschinot, C. Liu, and D. Krishnan, “Supervised contrastive learning”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [14] D. P. Kingma and J. Ba, *Adam: A method for stochastic optimization*, arXiv preprint arXiv:1412.6980, 2017. arXiv: 1412.6980 [cs.LG].
- [15] L. Hubert and P. Arabie, “Comparing partitions”, *Journal of Classification*, vol. 2, no. 1, pp. 193–218, 1985.
- [16] N. X. Vinh, J. Epps, and J. Bailey, “Information theoretic measures for clusterings comparison: Variants, properties, normalization and correction for chance”, *Journal of Machine Learning Research*, vol. 11, pp. 2837–2854, 2010.
- [17] A. Rosenberg and J. Hirschberg, “V-measure: A conditional entropy-based external cluster evaluation measure”, in *Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL)*, 2007, pp. 410–420.
- [18] H. W. Kuhn, “The Hungarian method for the assignment problem”, *Naval Research Logistics Quarterly*, vol. 2, no. 1–2, pp. 83–97, 1955. DOI: 10.1002/nav.3800020109
- [19] L. McInnes, J. Healy, and S. Astels, “hdbscan: Hierarchical density based clustering”, *The Journal of Open Source Software*, vol. 2, no. 11, p. 205, 2017. DOI: 10.21105/joss.00205
- [20] S. Bedoire, “Offline direction clustering of overlapping radar pulses from homogeneous emitters”, Host company Saab; DBSCAN-style clustering on pulse directions with simulated scenarios, Master’s thesis, KTH Royal Institute of Technology, Stockholm, Sweden, 2022.
- [21] V. Jakobsson, “Offline radar pulse track association with deinterleaving errors”, Host company SAAB; time-series clustering for track fragments under imperfect deinterleaving, Master’s thesis, KTH Royal Institute of Technology, Stockholm, Sweden, 2024.
- [22] K. Coleman, “Dataset drift in radar warning receivers: Out-of-distribution detection for radar emitter classification using an RNN-based deep ensemble”, UPTEC F 23041, 30 hp; simulator data from Saab AB and drift/OOD analysis for emitter classification, Master’s thesis, Uppsala University, Uppsala, Sweden, 2023.

## **APPENDIX A**

# **Additional Results**

Placeholder paragraph.

## **APPENDIX B**

# **Implementation Details**

Placeholder paragraph.



## APPENDIX C

# Notation Reference

Placeholder paragraph.